

model\_result\_008

Yi Huang, Yuchen Zhang, Shun Xie

2023-05-09

## Create and Save data

```
#load data
load("data/recovery.RData")

#get data from 2 team members uni
set.seed(3554)
dat1 <- dat[sample(1:10000, 2000),]
set.seed(4437)
dat2 <- dat[sample(1:10000, 2000),]

#combine the data and discard duplicates
dat_temp <- rbind(dat1, dat2)
dat <- dat_temp[!duplicated(dat_temp$id),]

#dat is now the final data and it has the length as following:
length(dat$id)
```

```
## [1] 3593
```

```
data =
  dat %>%
  mutate(study = factor(dat$study),
         gender = factor(dat$gender),
         hypertension = factor(dat$hypertension),
         diabetes = factor(dat$diabetes),
         vaccine = factor(dat$vaccine),
         severity = factor(dat$severity)) %>%
  dplyr::select(-id)

#save data
save(data, file = "data/covid_recovery.RData")
head(data)
```

```
##      age gender race smoking height weight  bmi hypertension diabetes SBP LDL
## 2989  63      1    1      0  174.8   83.6 27.4           1         0  137 108
## 620   60      1    1      0  170.0   80.7 27.9           0         0  128  80
## 9226  64      0    1      0  167.5   89.1 31.7           1         0  137 126
## 4098  69      0    1      2  173.0   72.5 24.2           1         0  135 149
## 8349  62      0    1      0  172.6   82.4 27.7           1         0  138 136
## 184   60      1    1      0  166.5   82.9 29.9           0         0  128 118
##      vaccine severity study recovery_time
## 2989         1         0    B           14
## 620          1         0    A           36
## 9226         0         0    C           50
## 4098         0         0    B           65
## 8349         1         0    C           34
## 184          1         0    A           32
```

## Load data, data partition

```
# load data
load("data/covid_recovery.Rdata")

#Split data into 70-300, using the third member's uni
```

```

set.seed(2337)
indexTrain <- createDataPartition(y = data$recovery_time, p = 0.7, list = FALSE)

# training data
train_data <- data[indexTrain,]
# matrix of predictors
train_x <- model.matrix(recovery_time~.,train_data)[,-1]
# vector of response
train_y <- train_data$recovery_time

# test data
test_data <- data[-indexTrain,]
# matrix of predictors
test_x <- model.matrix(recovery_time~.,test_data)[,-1]
# vector of response
test_y <- test_data$recovery_time

# 10-fold cv
ctrl <- trainControl(method = "cv",
                     number = 10)


```

## EDA on training data

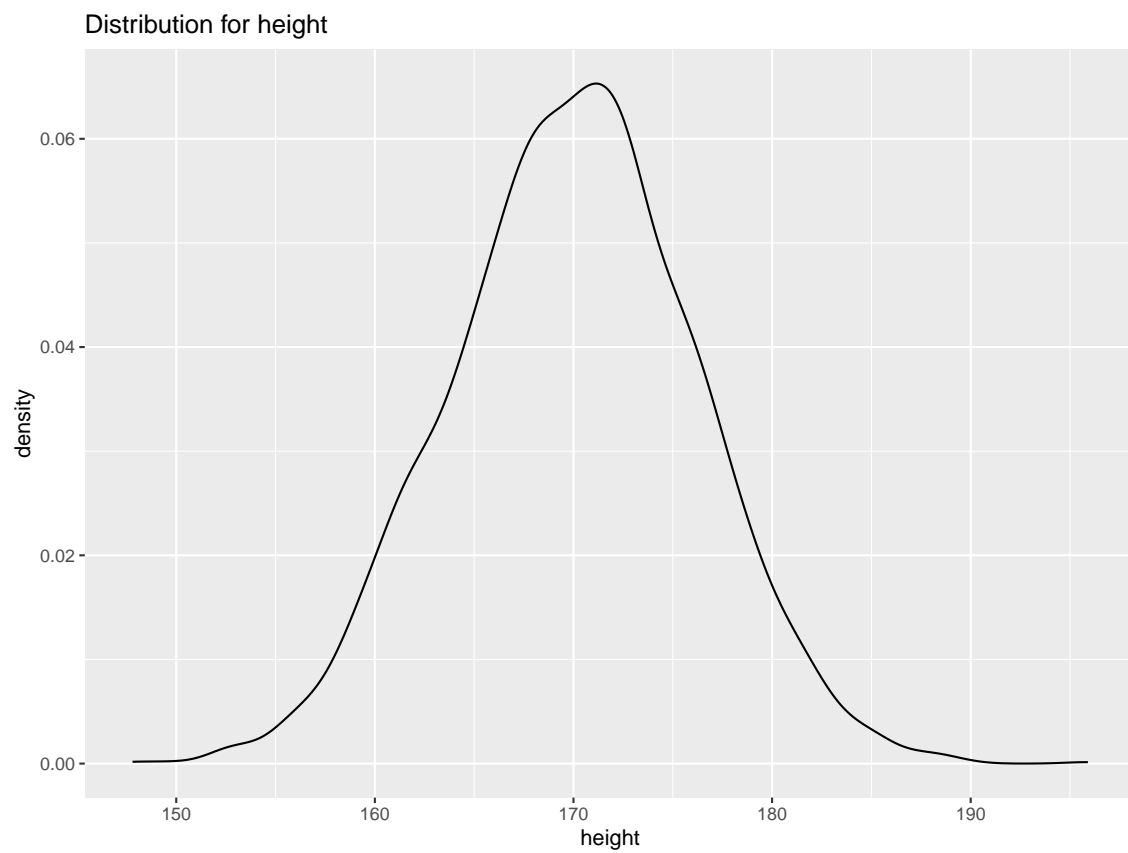
```

#continuous
ggplot(train_data, aes(x = age))+geom_density()+labs(title="Distribution for age")

```




```
ggplot(train_data, aes(x = height))+geom_density()+labs(title="Distribution for height")
```

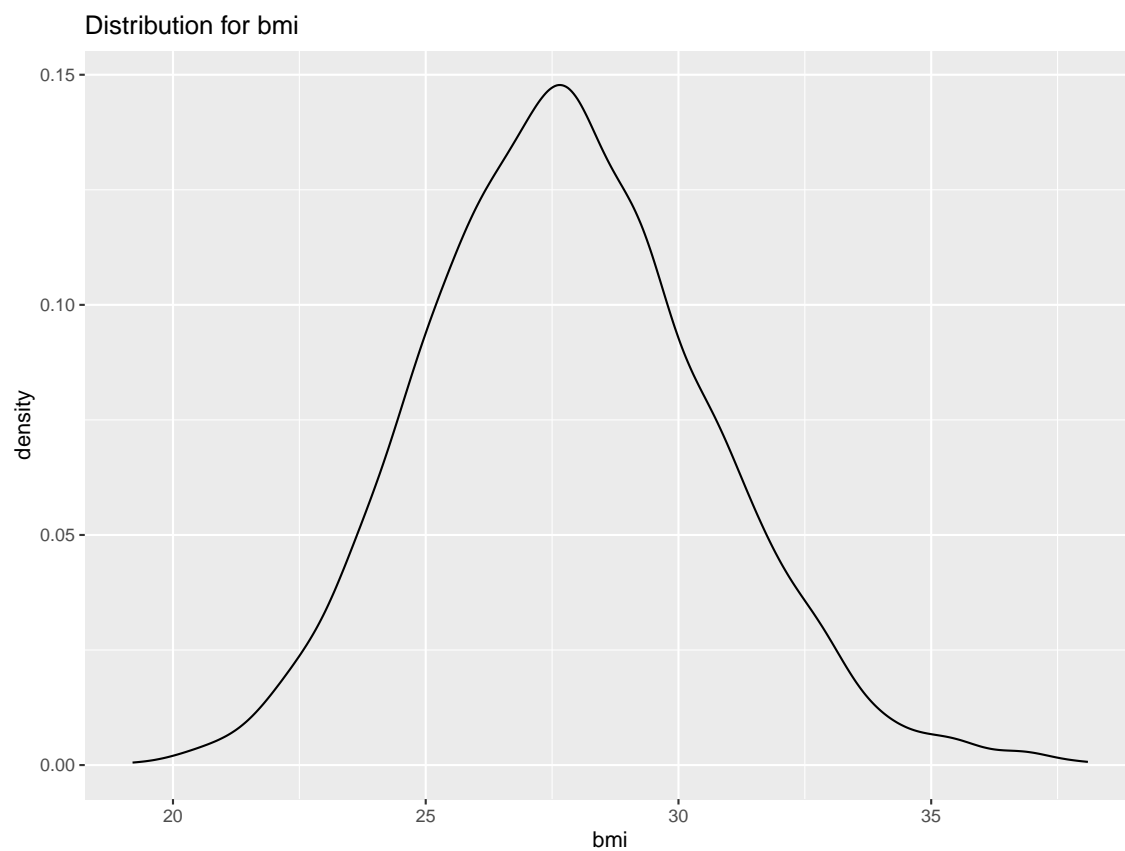


```
ggplot(train_data, aes(x = weight))+geom_density()+labs(title="Distribution for weight")
```

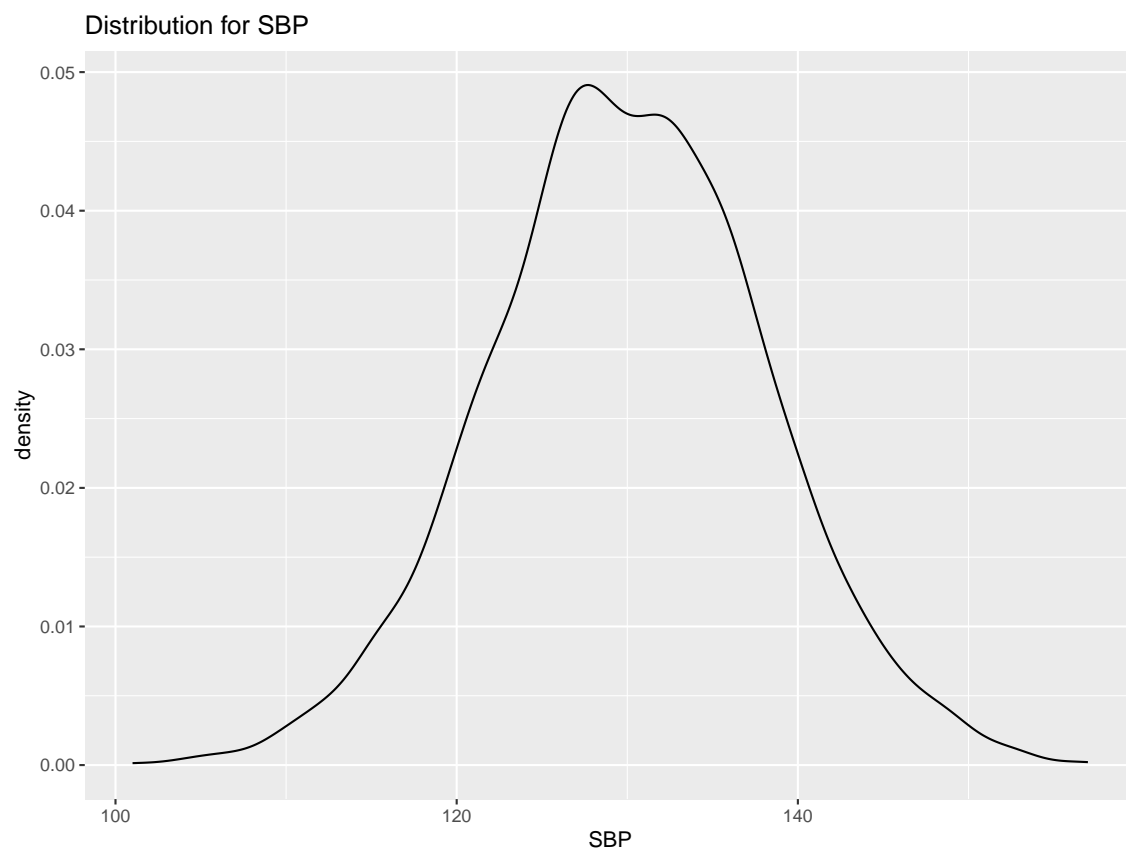
Distribution for weight



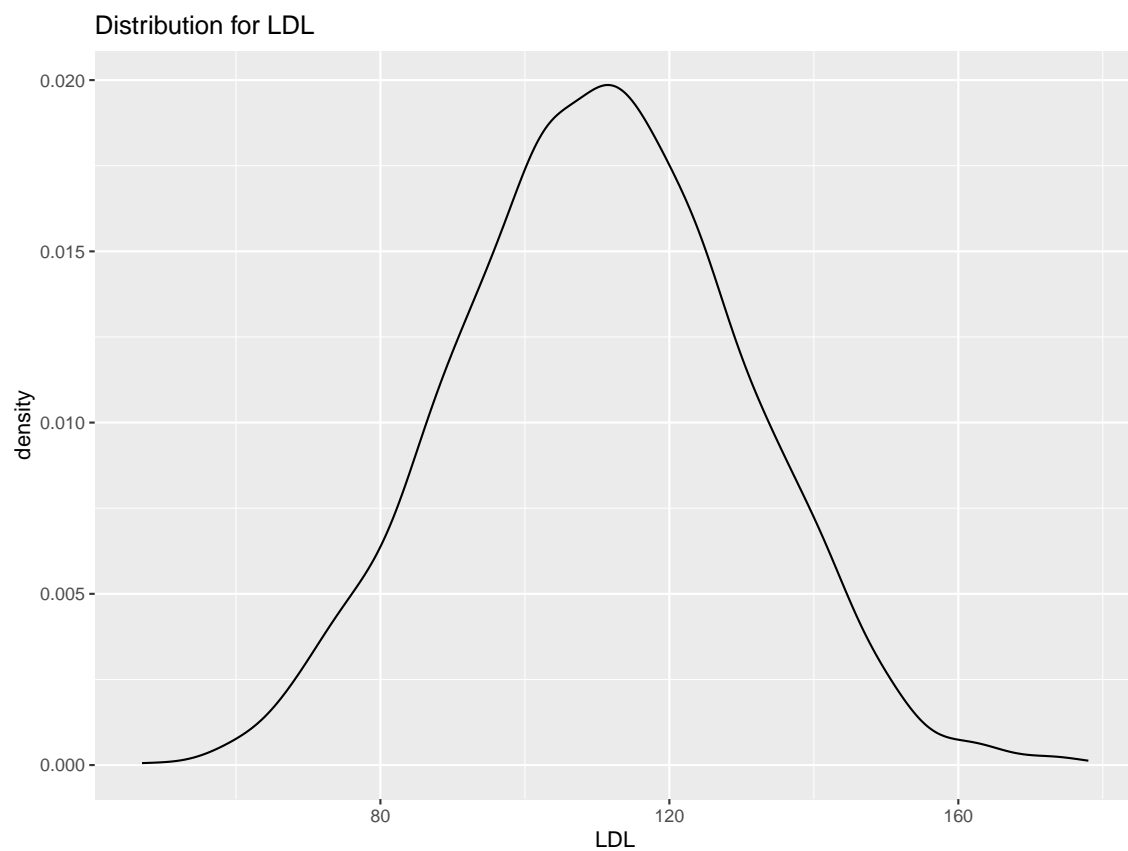
```
ggplot(train_data, aes(x = bmi))+geom_density()+labs(title="Distribution for bmi")
```



```
ggplot(train_data, aes(x = SBP))+geom_density()+labs(title="Distribution for SBP")
```



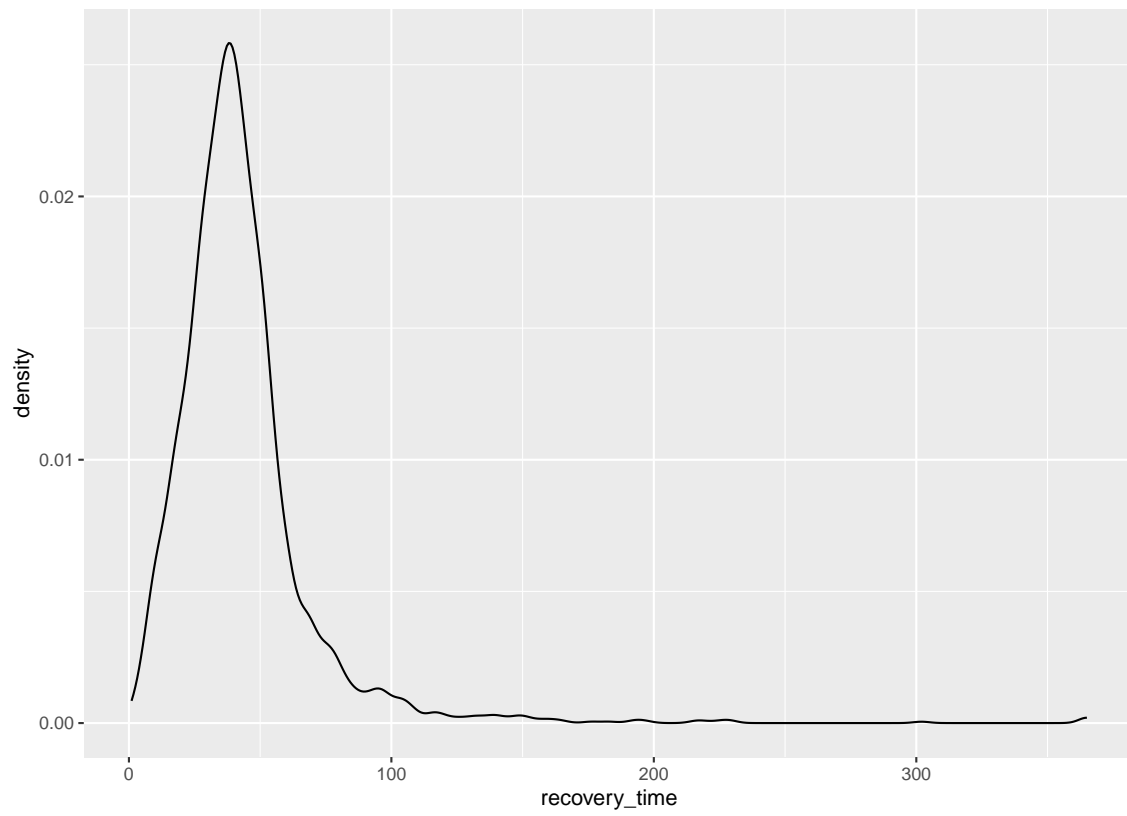
```
ggplot(train_data, aes(x = LDL))+geom_density()+labs(title="Distribution for LDL")
```




```
ggplot(train_data, aes(x = recovery_time))+geom_density()+labs(title="Distribution for recovery_time")
```



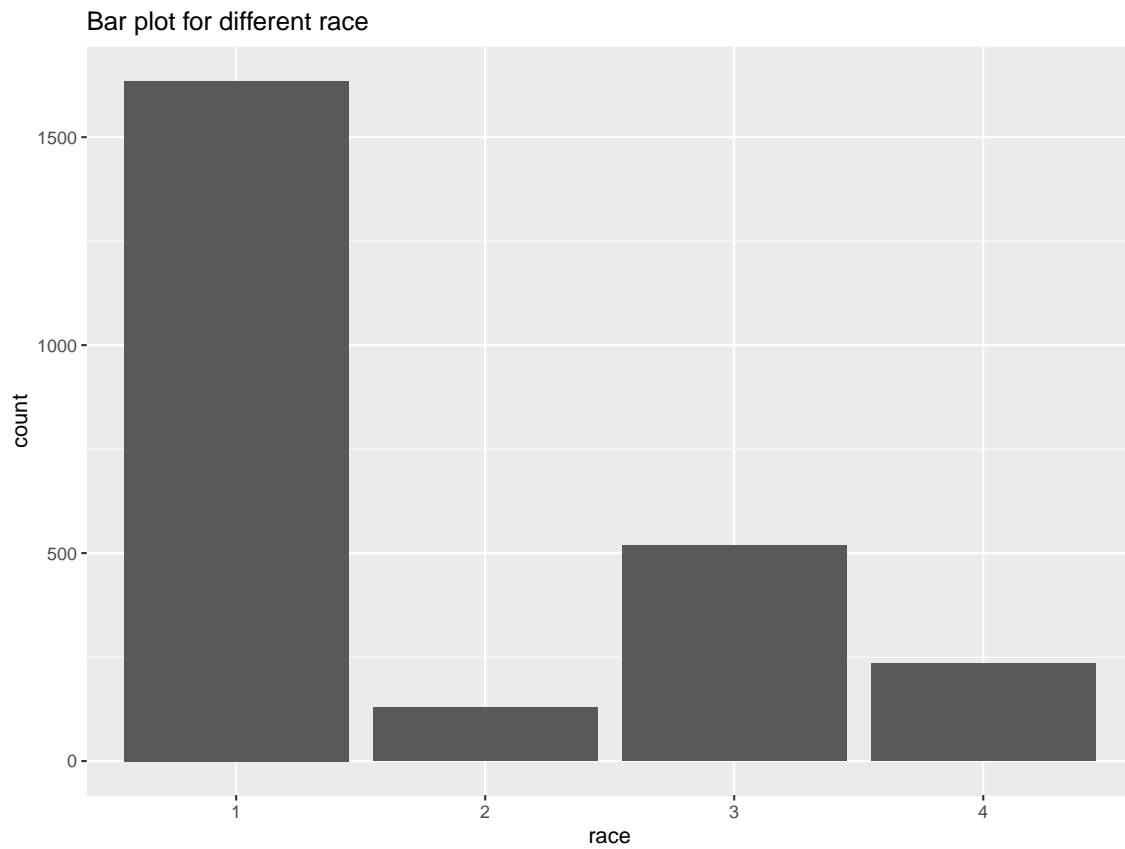
Distribution for recovery\_time



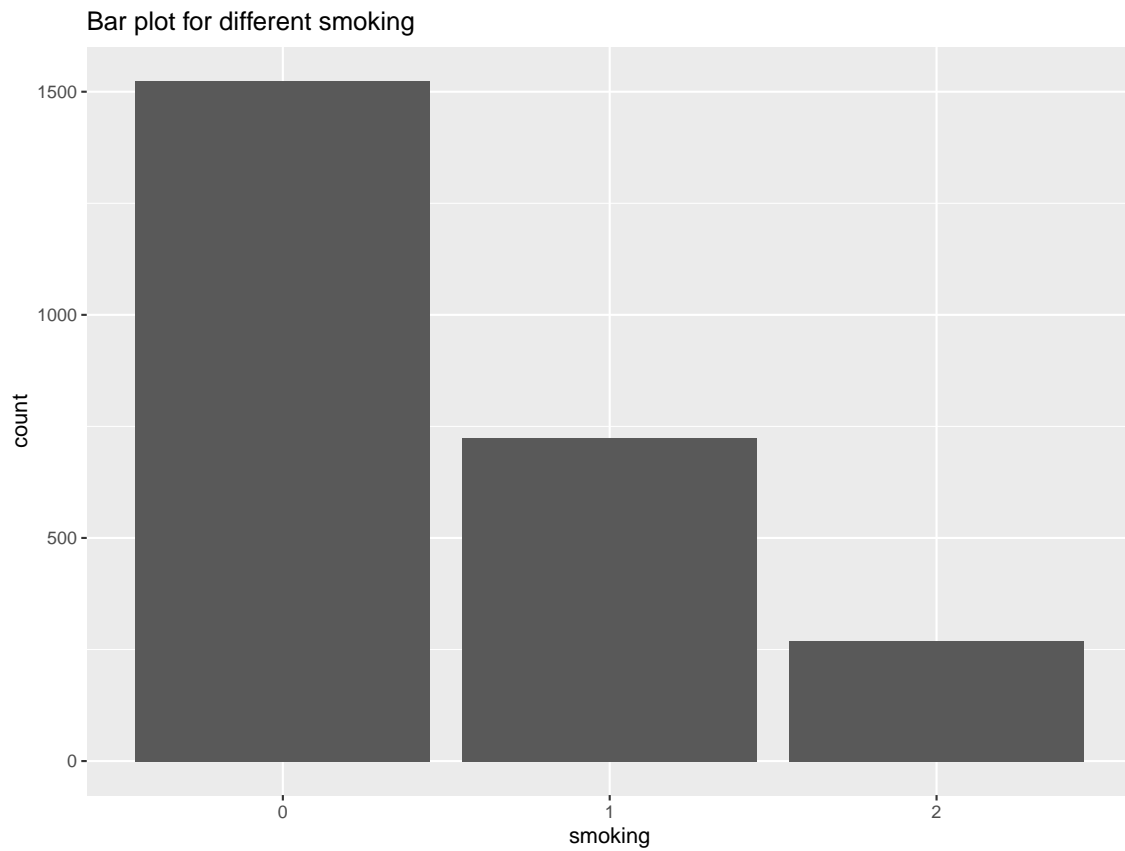
```
#discrete  
ggplot(train_data) + geom_bar(aes(x=gender))+labs(title="Bar plot for different gender")
```



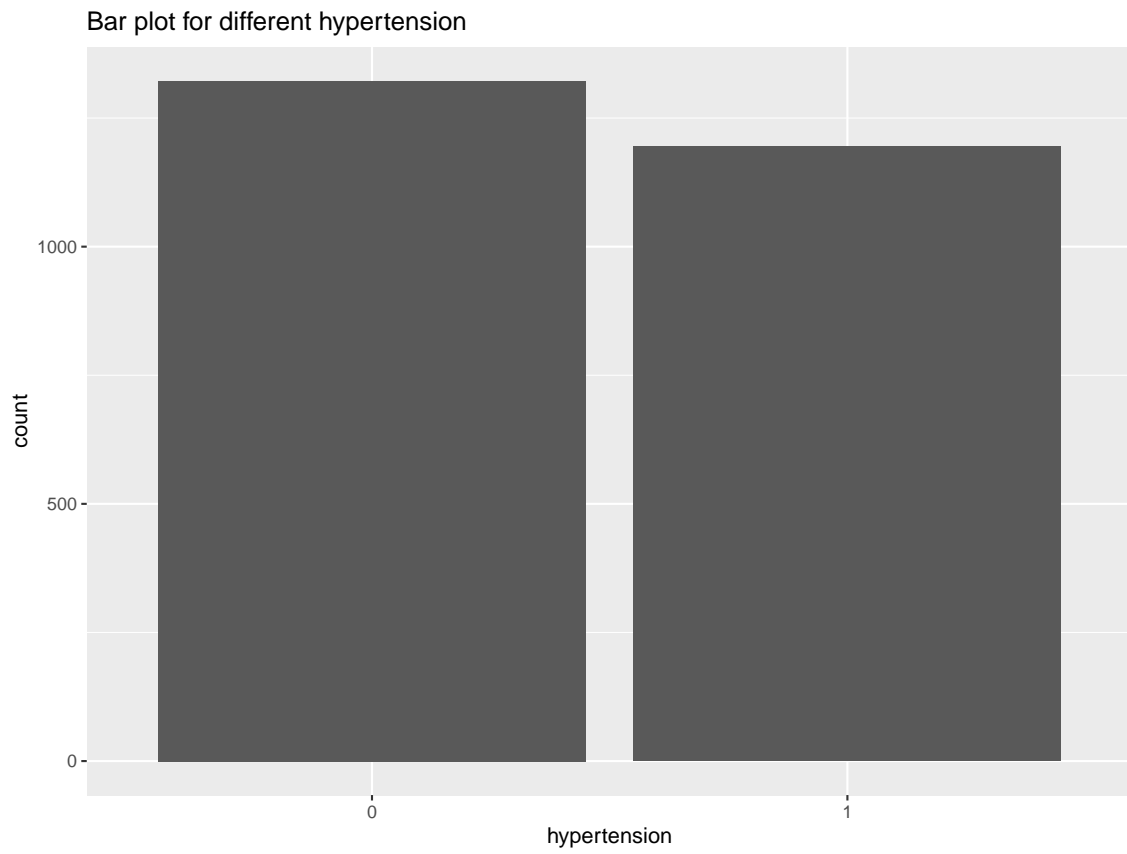
```
ggplot(train_data) + geom_bar(aes(x=gender))+labs(title="Bar plot for different gender")
```



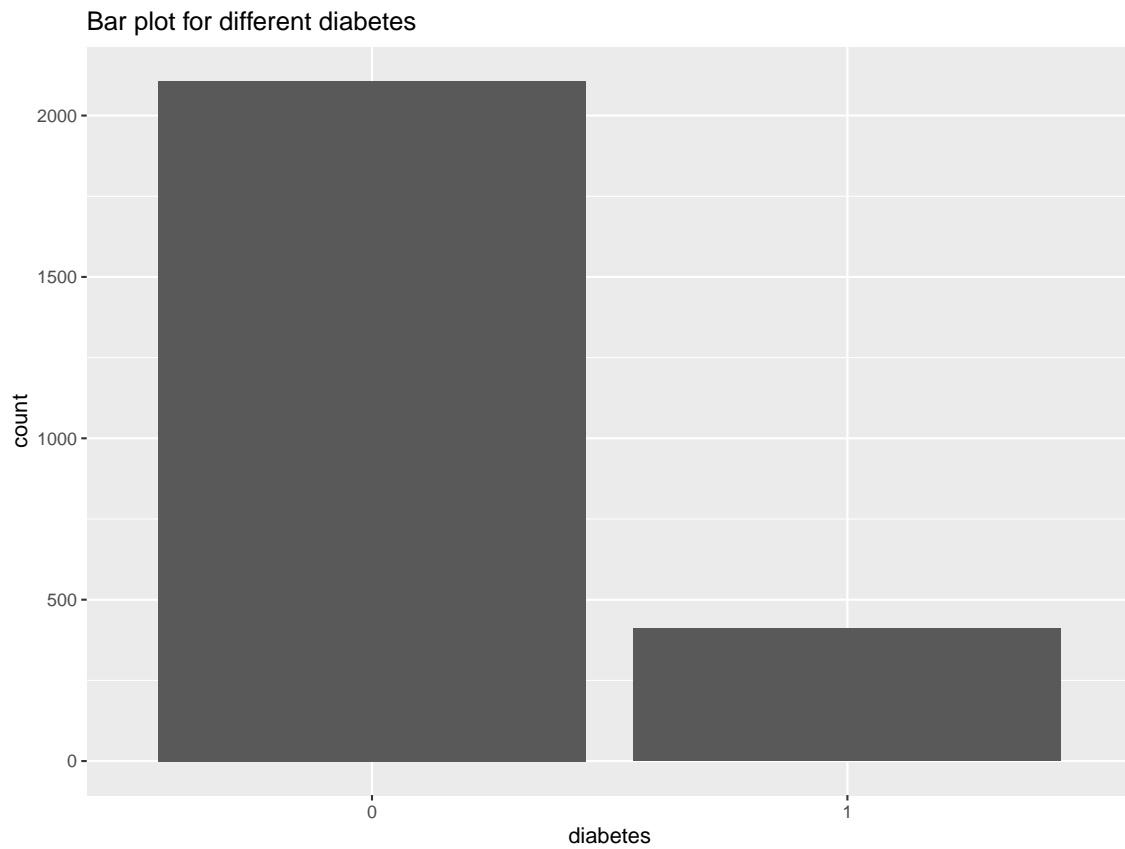
```
ggplot(train_data) + geom_bar(aes(x=smoking))+labs(title="Bar plot for different smoking")
```



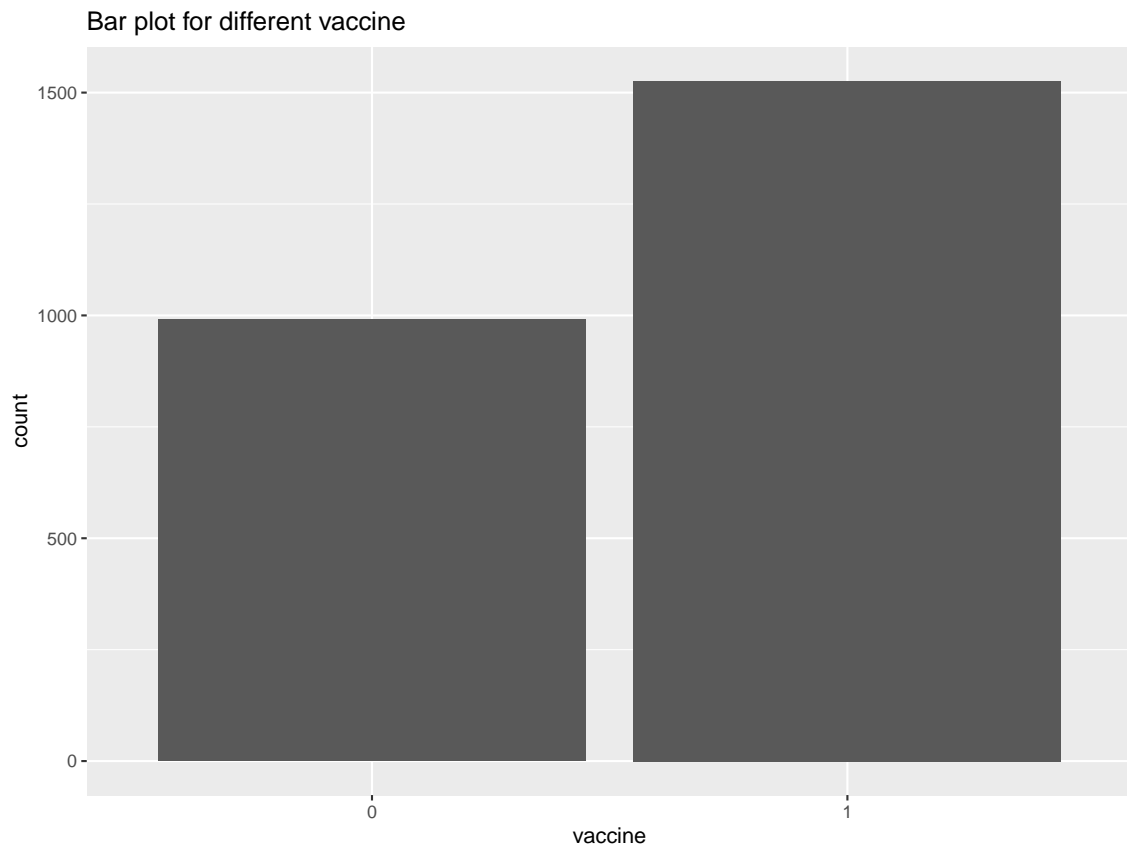
```
ggplot(train_data) + geom_bar(aes(x=hypertension))+labs(title="Bar plot for different hypertension")
```



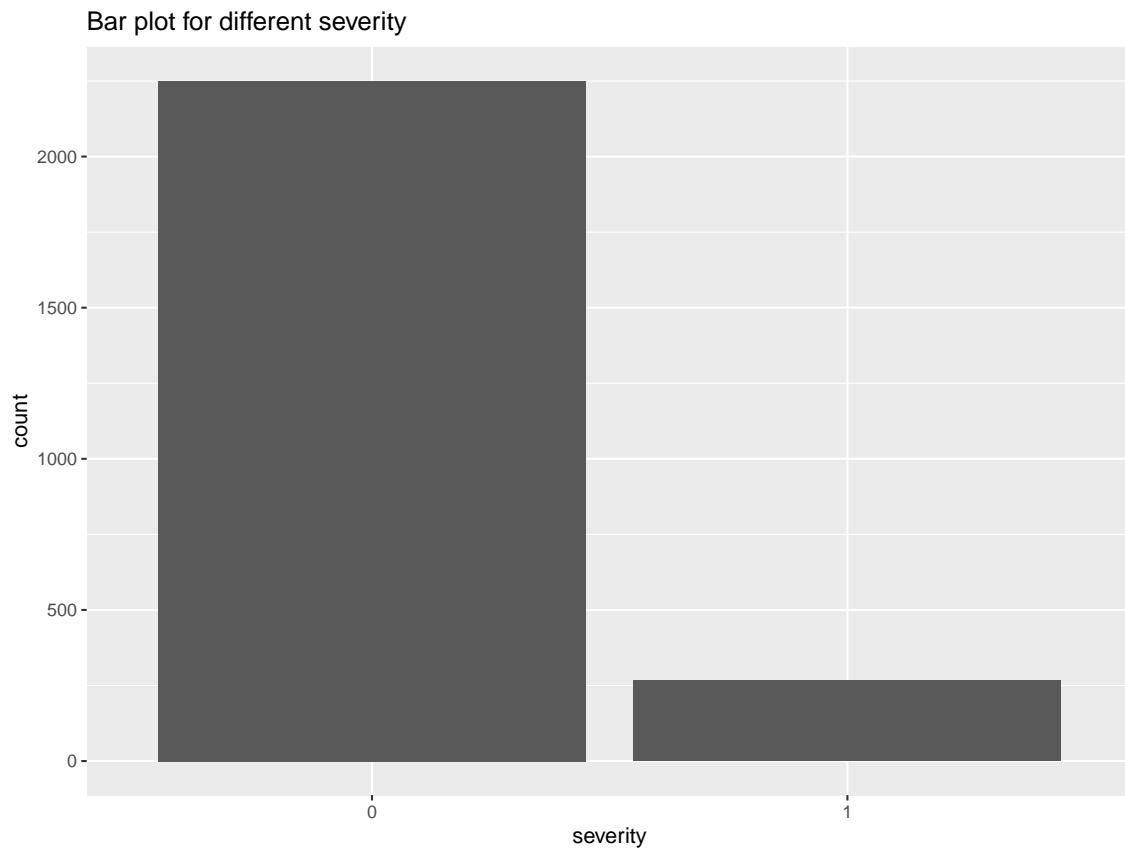
```
ggplot(train_data) + geom_bar(aes(x=diabetes))+labs(title="Bar plot for different diabetes")
```



```
ggplot(train_data) + geom_bar(aes(x=vaccine))+labs(title="Bar plot for different vaccine")
```

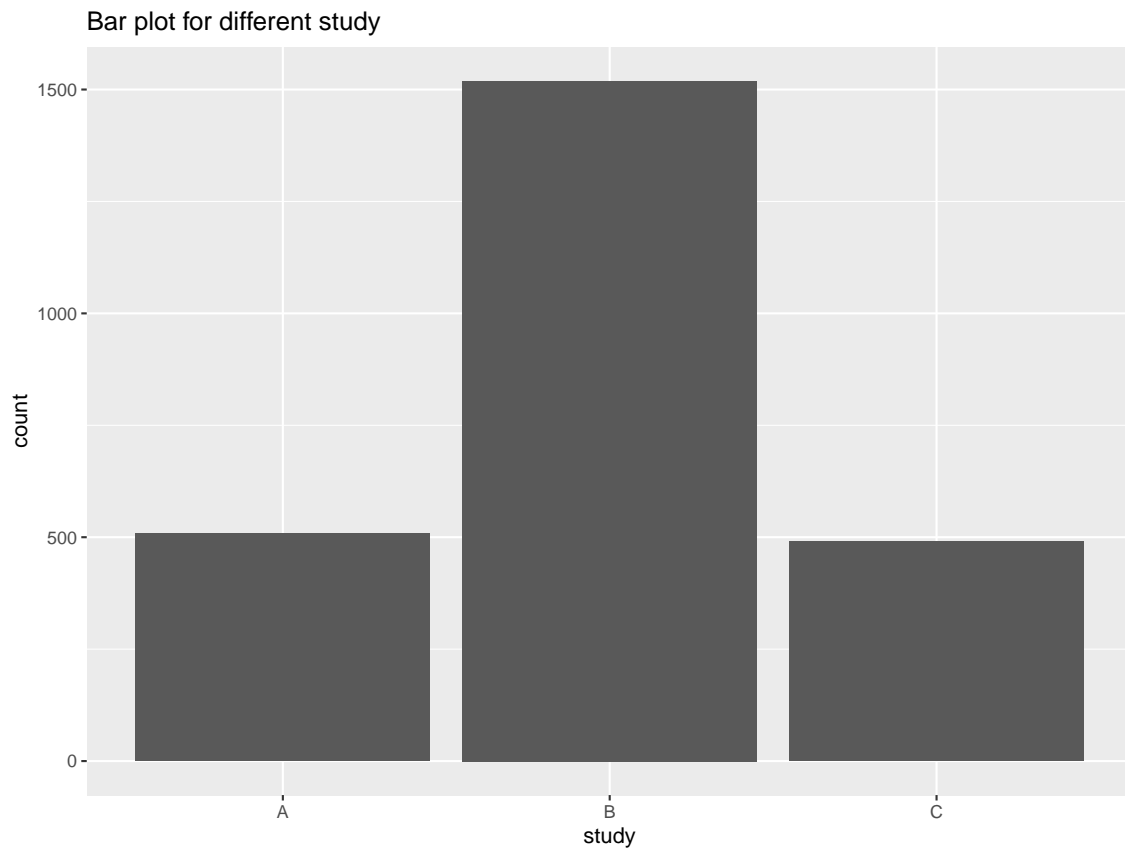


```
ggplot(train_data) + geom_bar(aes(x=severity))+labs(title="Bar plot for different severity")
```



```
ggplot(train_data) + geom_bar(aes(x=study))+labs(title="Bar plot for different study")
```



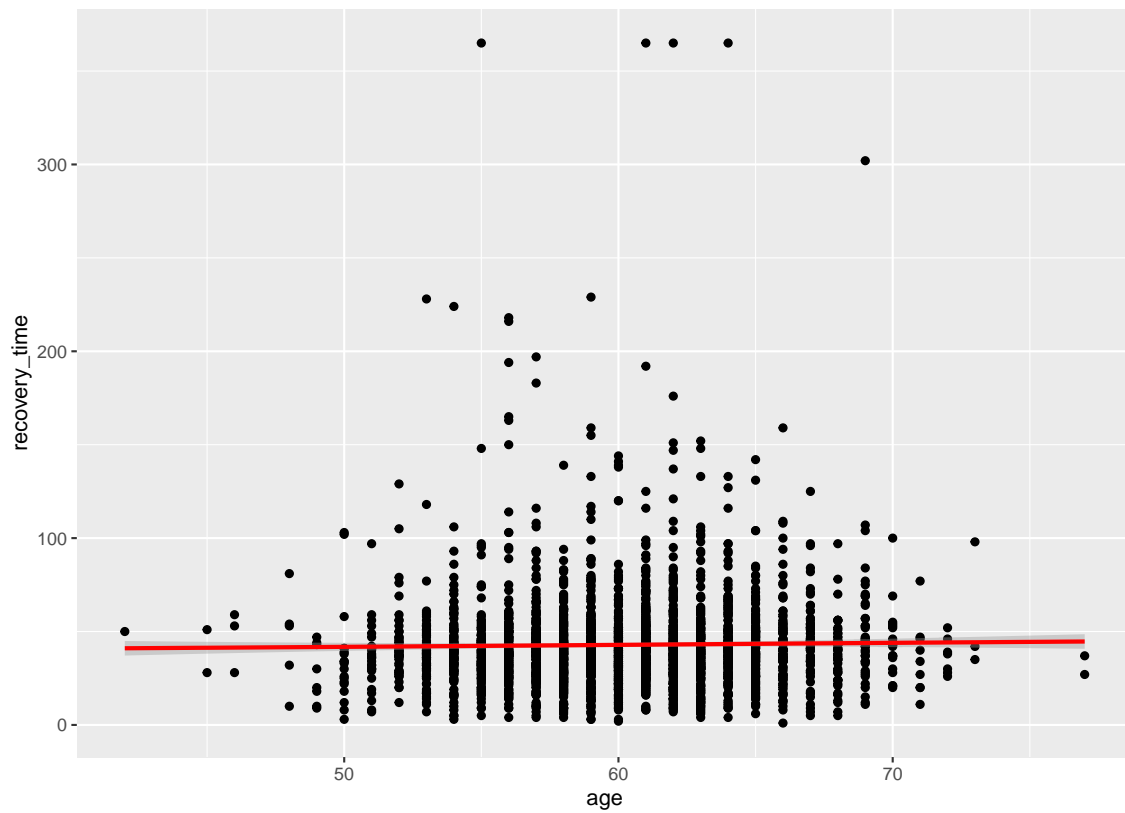


```
num_df <-  
  train_data %>%  
  dplyr::select(age, height, weight, bmi, SBP, LDL, recovery_time)  
  
# calculate the correlations  
res <- cor(num_df, use="complete.obs")  
  
corrplot(res, type = "upper", order = "hclust",  
          tl.col = "black", tl.srt = 45)
```



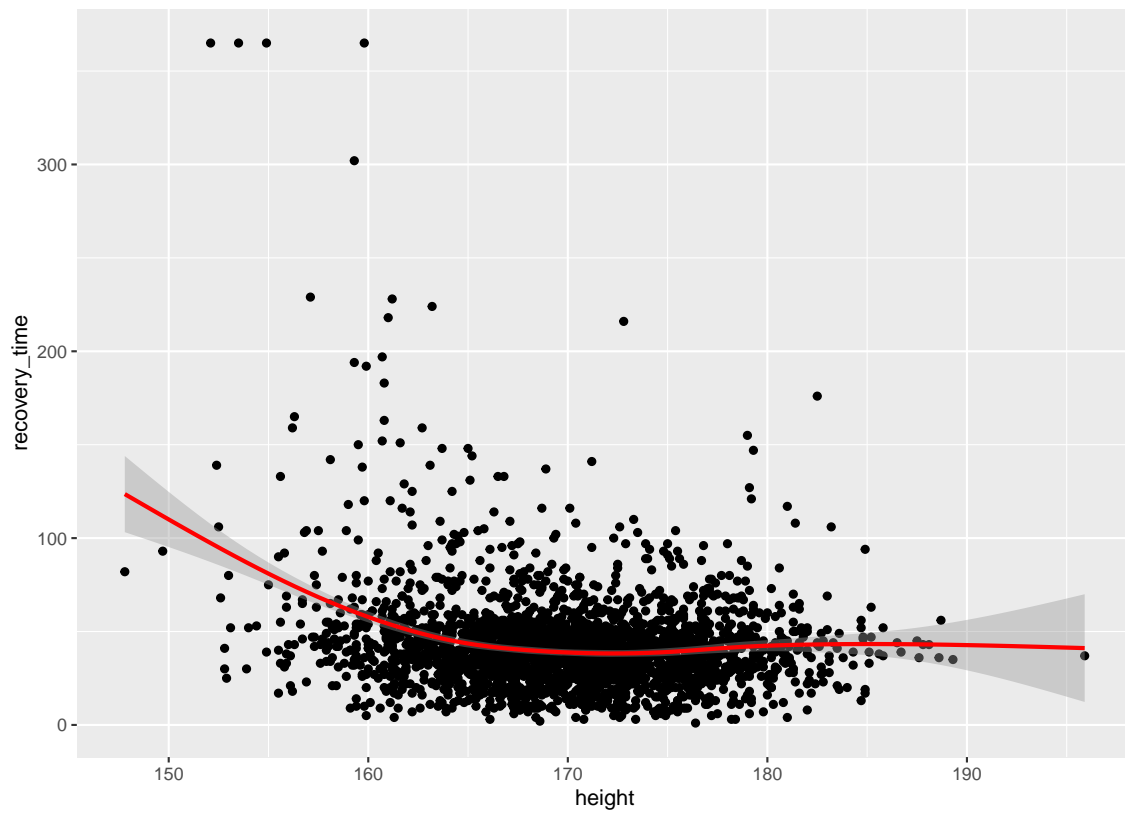
```
ggplot(train_data, aes(x = age, y = recovery_time))+geom_point()+geom_smooth(method = 'gam', se = TRUE,
## `geom_smooth()` using formula = 'y ~ s(x, bs = "cs")'
```

Recovery time against age



```
ggplot(train_data, aes(x = height, y = recovery_time))+geom_point()+geom_smooth(method = 'gam', se = TRUE)
## `geom_smooth()` using formula = 'y ~ s(x, bs = "cs")'
```

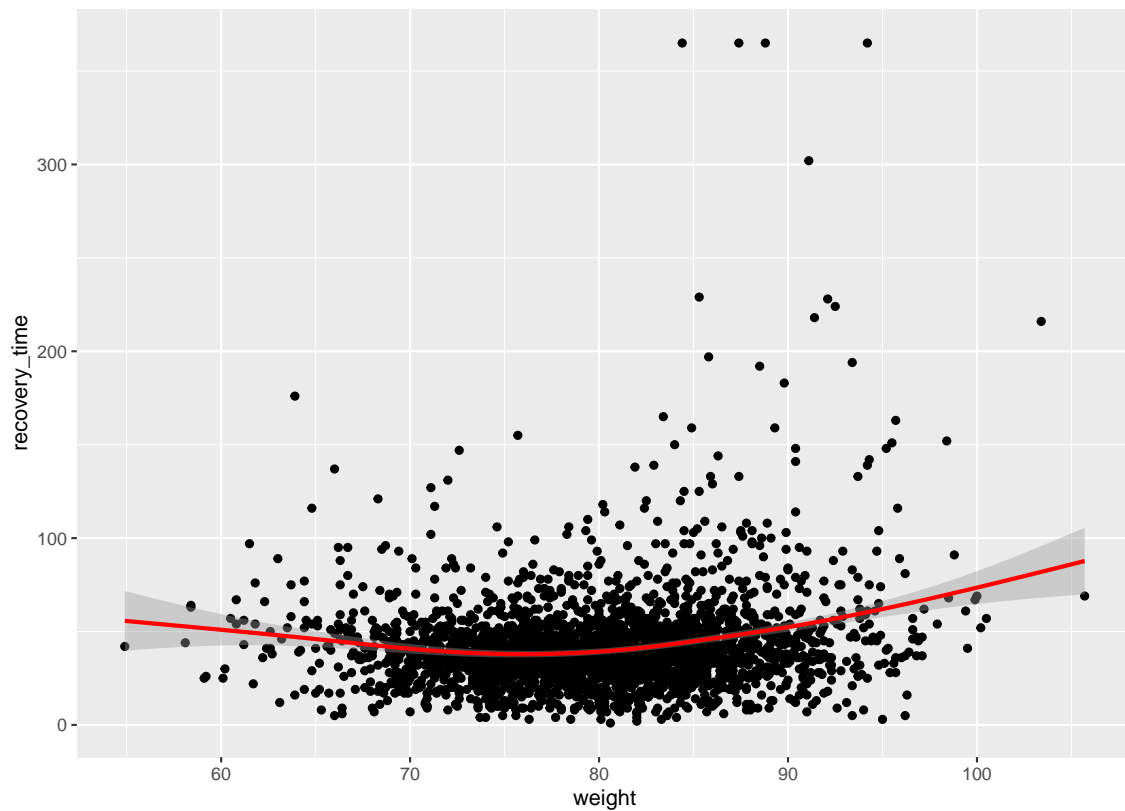
Recovery time against height



```
ggplot(train_data, aes(x = weight, y = recovery_time))+geom_point()+geom_smooth(method = 'gam', se = TRUE)

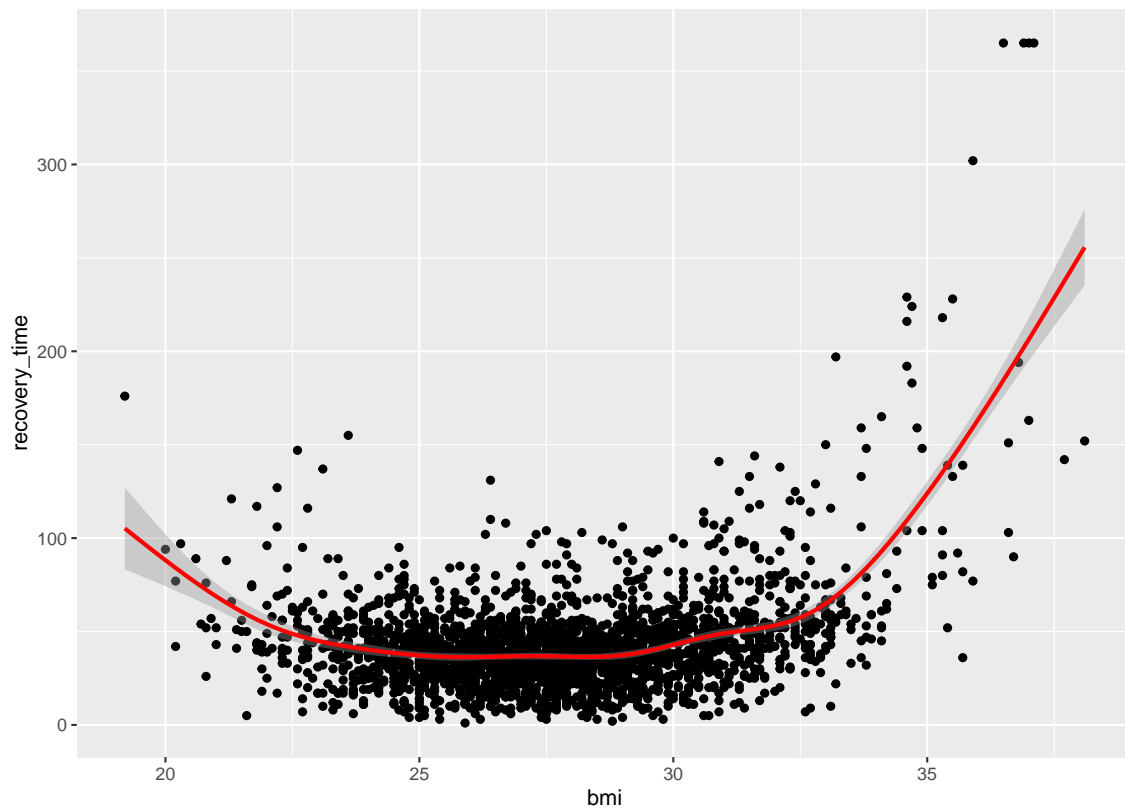
## `geom_smooth()` using formula = 'y ~ s(x, bs = "cs")'
```

Recovery time against weight



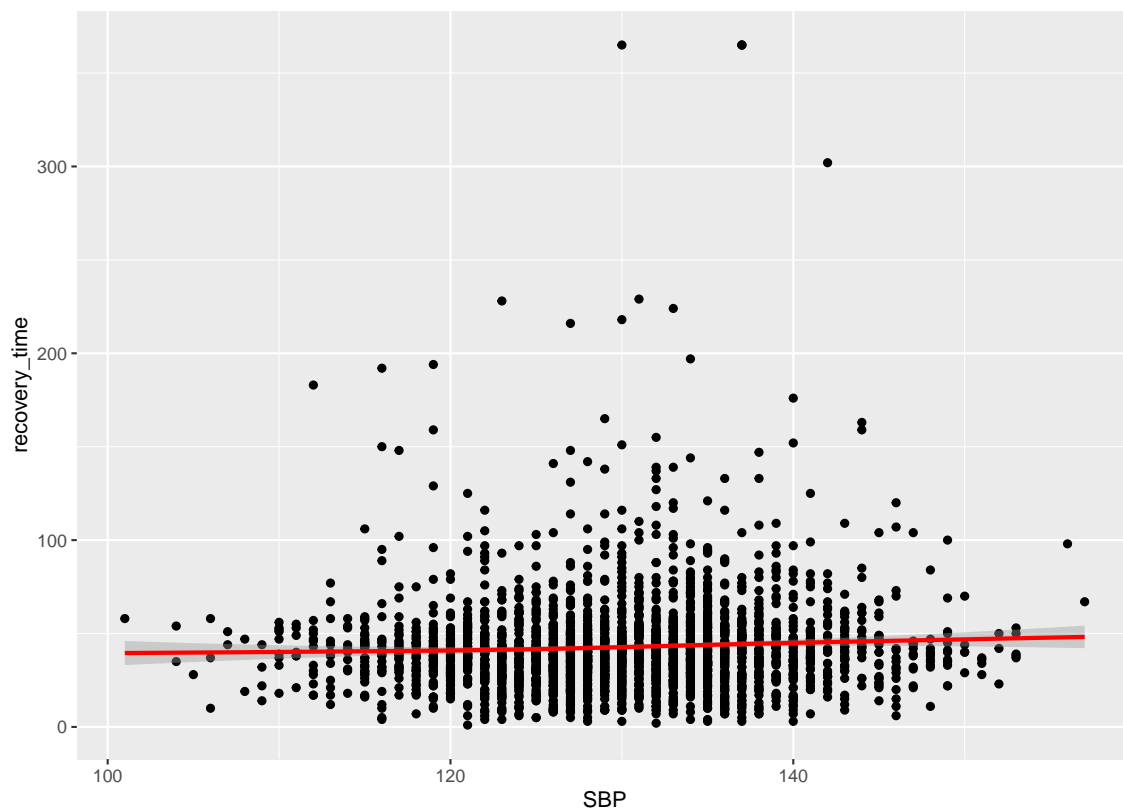
```
ggplot(train_data, aes(x = bmi, y = recovery_time))+geom_point()+geom_smooth(method = 'gam', se = TRUE,  
## `geom_smooth()` using formula = 'y ~ s(x, bs = "cs")'
```

Recovery time against bmi




```
ggplot(train_data, aes(x = BMI, y = recovery_time))+geom_point()+geom_smooth(method = 'gam', se = TRUE,  
## `geom_smooth()` using formula = 'y ~ s(x, bs = "cs")'
```

Recovery time against SBP



```
ggplot(train_data, aes(x = LDL, y = recovery_time))+geom_point()+geom_smooth(method = 'gam', se = TRUE,  
## `geom_smooth()` using formula = 'y ~ s(x, bs = "cs")'
```


Recovery time against LDL



```
ggplot(data, aes(x = as.factor(gender), y = recovery_time))+geom_boxplot()+labs(title="Box plot for Rec
```

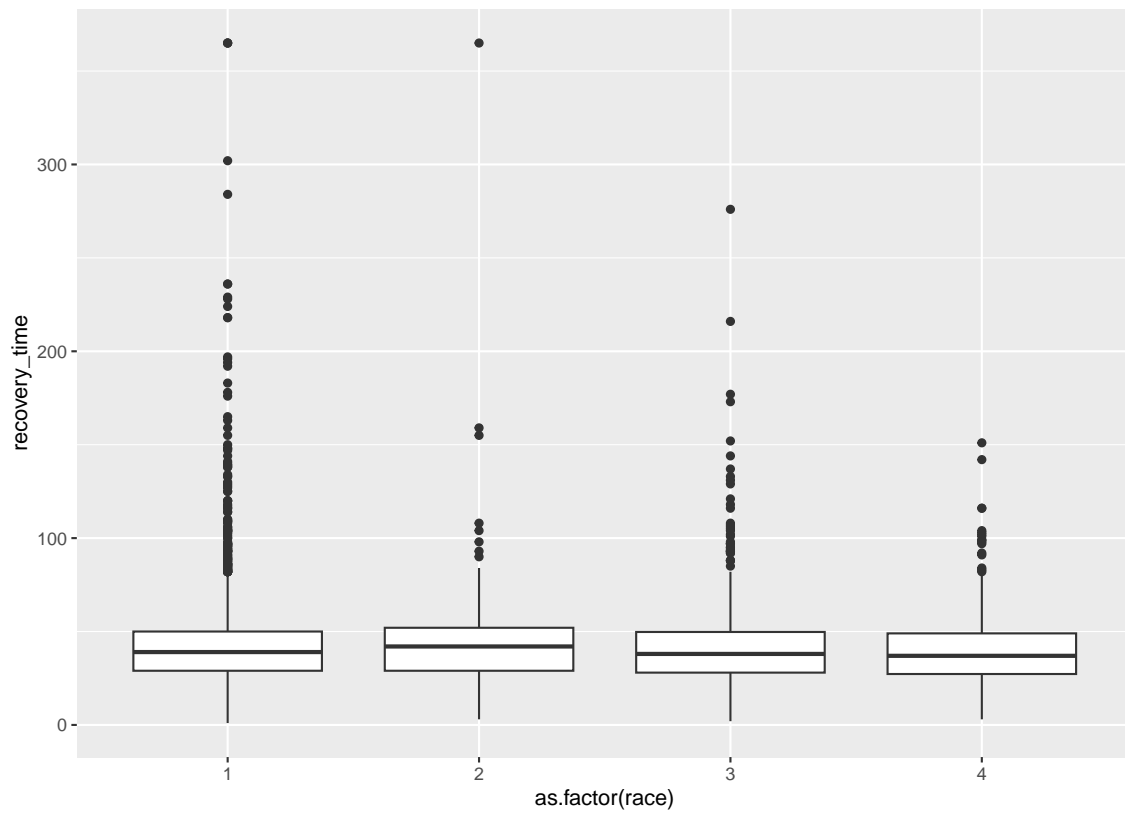


Box plot for Recovery time with different gender



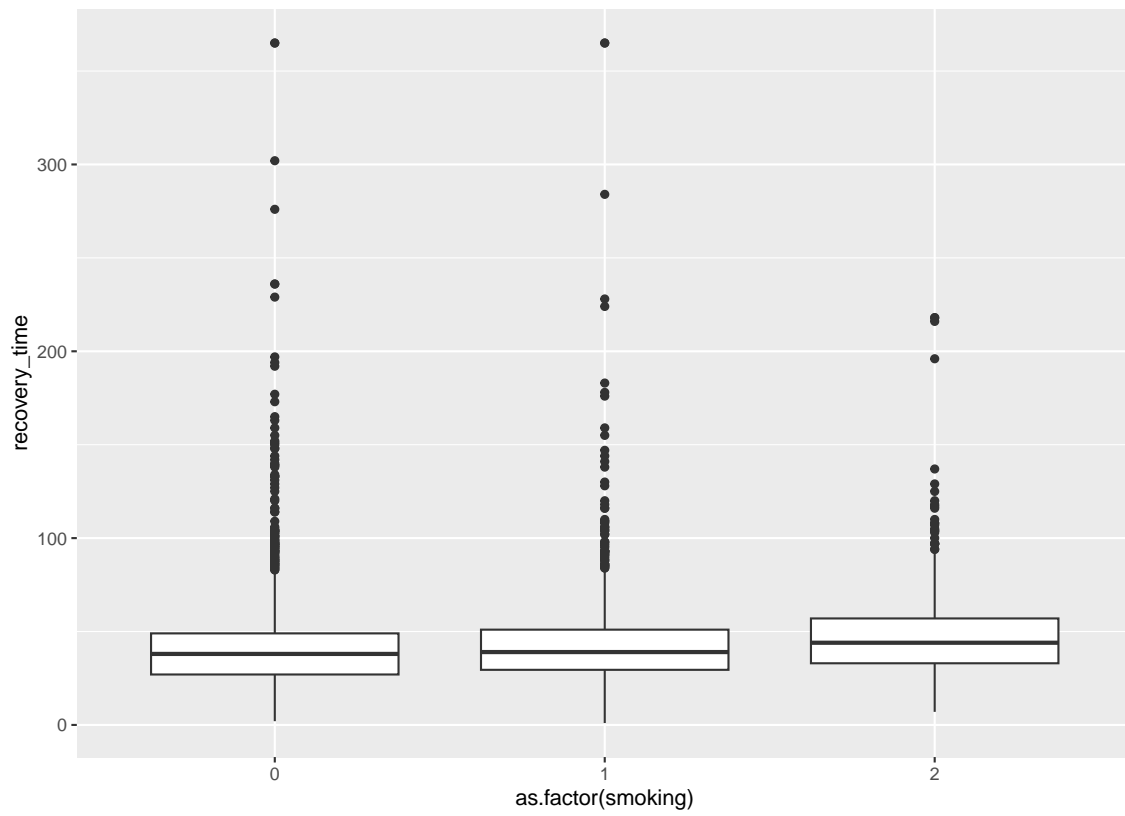
```
ggplot(data, aes(x = as.factor(race), y = recovery_time))+geom_boxplot()+labs(title="Box plot for Recov
```

Box plot for Recovery time with different race



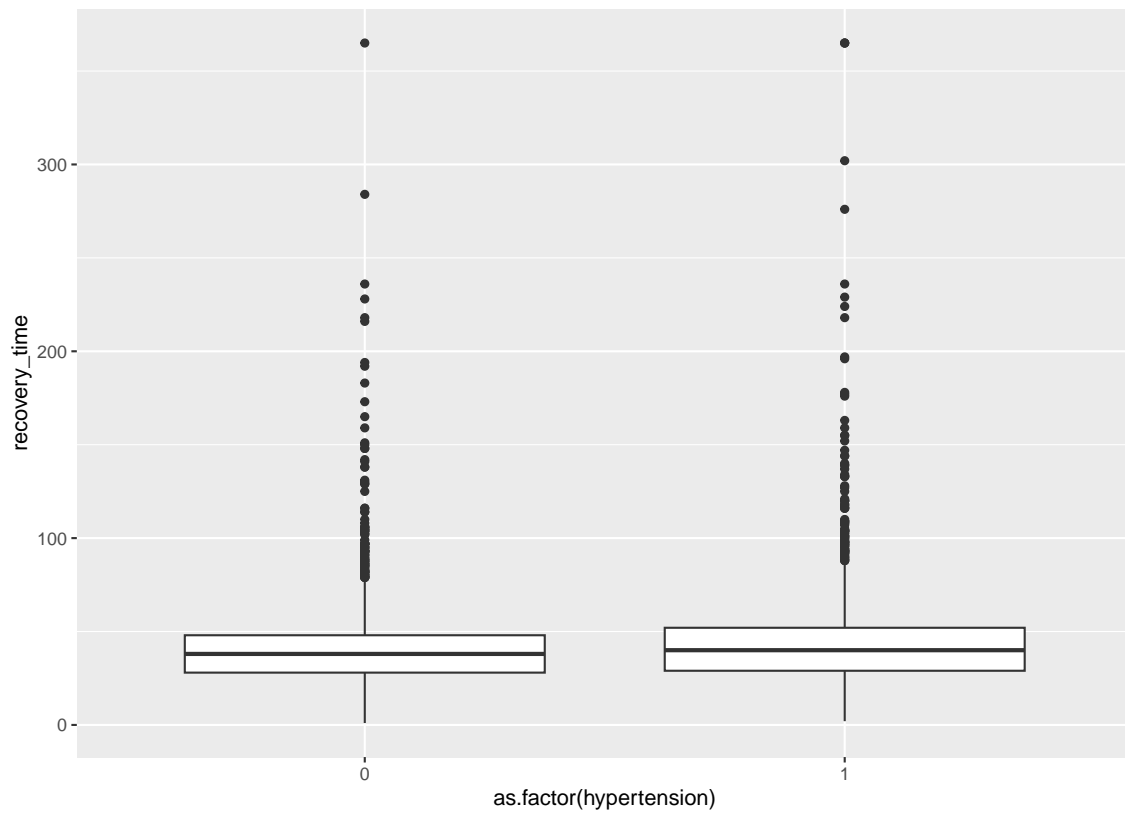
```
ggplot(data, aes(x = as.factor(smoking), y = recovery_time))+geom_boxplot()+labs(title="Box plot for Re
```

Box plot for Recovery time with different smoking



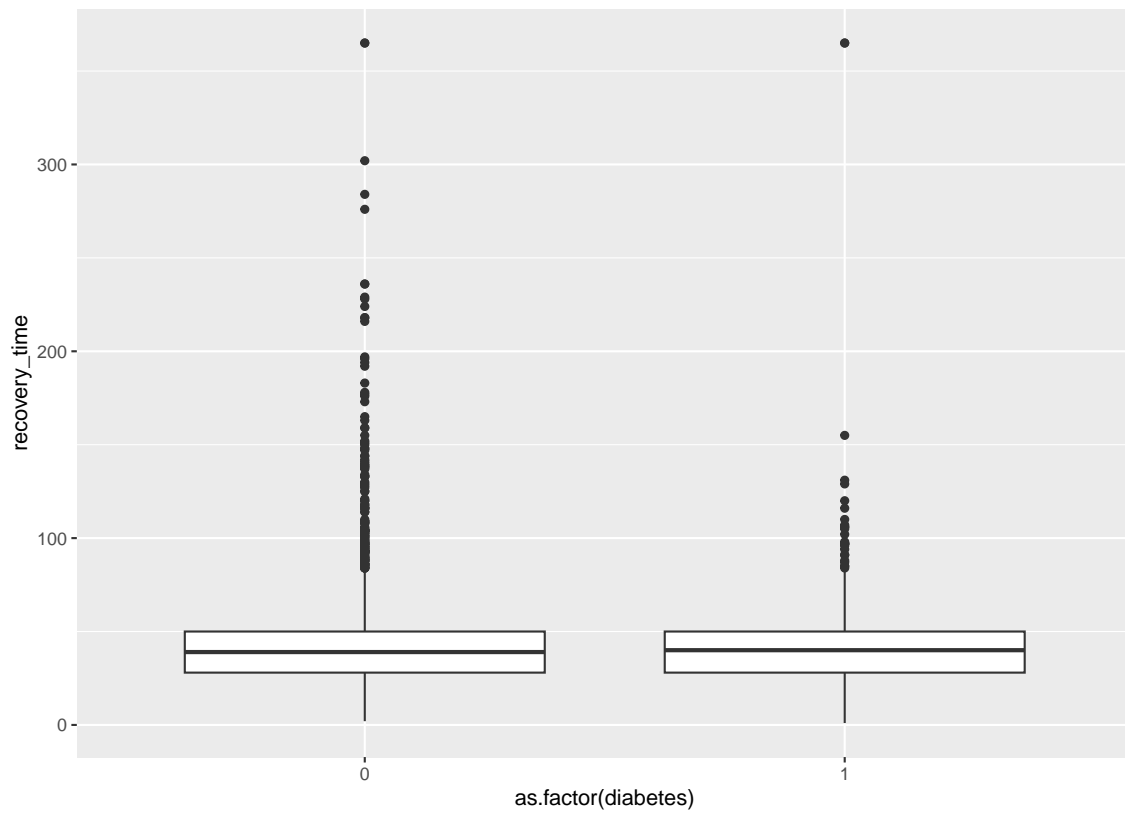
```
ggplot(data, aes(x = as.factor(hypertension), y = recovery_time))+geom_boxplot()+labs(title="Box plot f
```

Box plot for Recovery time with different hypertension




```
ggplot(data, aes(x = as.factor(diabetes), y = recovery_time))+geom_boxplot()+labs(title="Box plot for R
```

Box plot for Recovery time with different diabetes



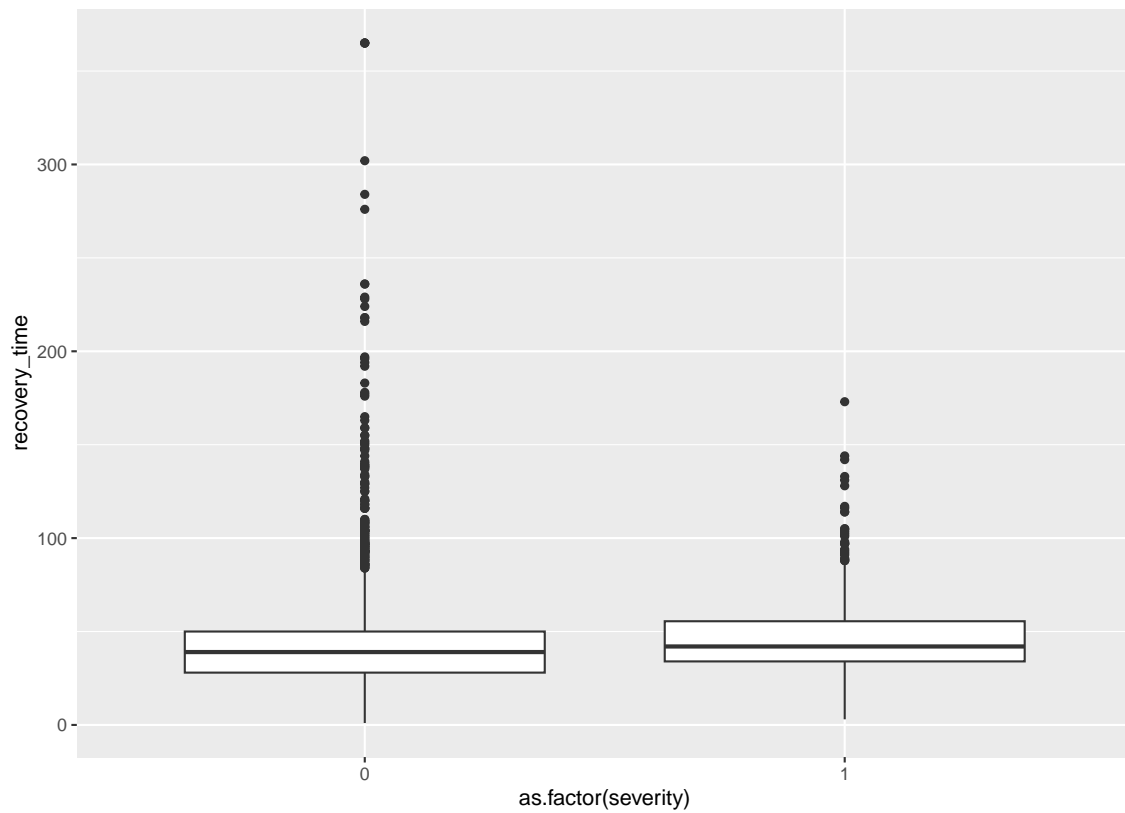
```
ggplot(data, aes(x = as.factor(vaccine), y = recovery_time))+geom_boxplot()+labs(title="Box plot for Re
```

Box plot for Recovery time with different vaccine

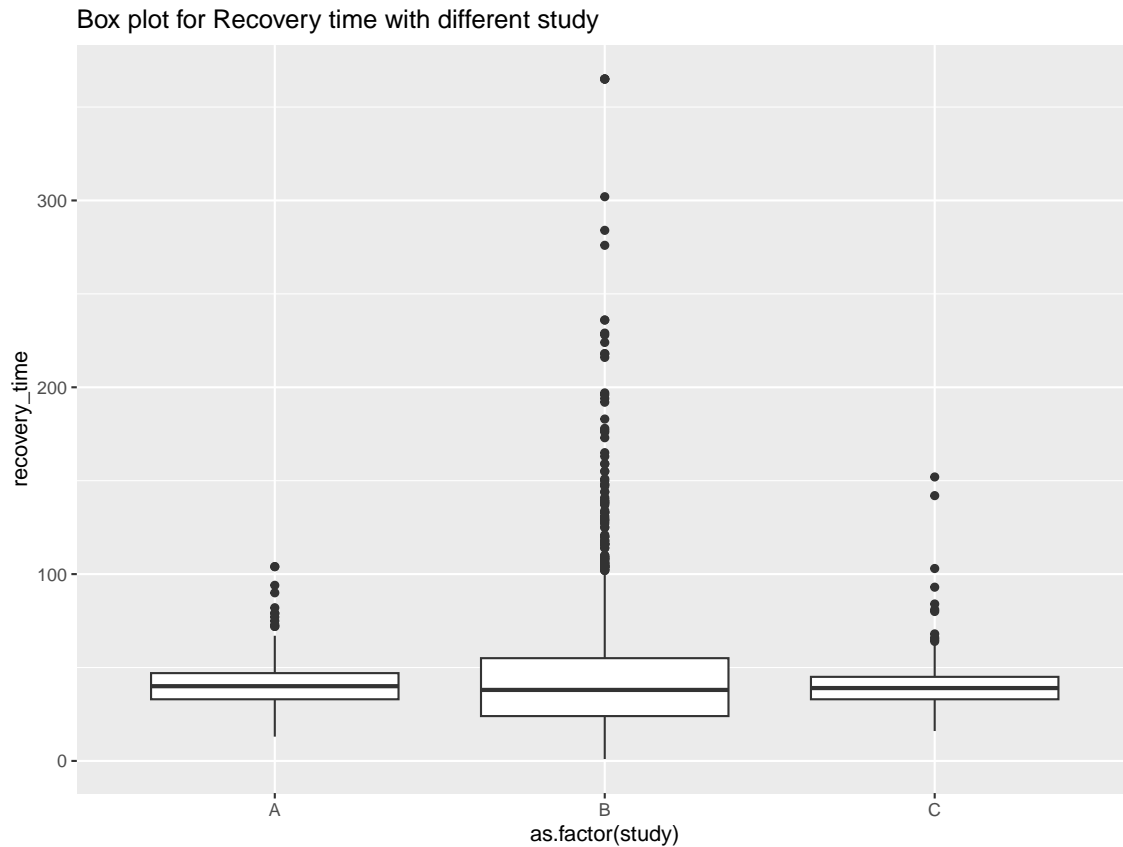


```
ggplot(data, aes(x = as.factor(severity), y = recovery_time))+geom_boxplot()+labs(title="Box plot for R
```

Box plot for Recovery time with different severity



```
ggplot(data, aes(x = as.factor(severity), y = recovery_time))+geom_boxplot()+labs(title="Box plot for Recovery time with different severity")
```



## Regression

### primary analysis (Regression)

- Linear regression
- K-Nearest Neighbors (KNN)
- Elastic net
- Partial least squares (PLS)
- Generalized Additive Model (GAM)
- Multivariate Adaptive Regression Splines (MARS)
- Boosting
- Random forest

### Linear regression

```
## fit linear model on train data
set.seed(8)
linear_model <- train(train_x,
                      train_y,
                      method = "lm",
                      trControl = ctrl)
summary(linear_model)

##
## Call:
## lm(formula = .outcome ~ ., data = dat)
##
```



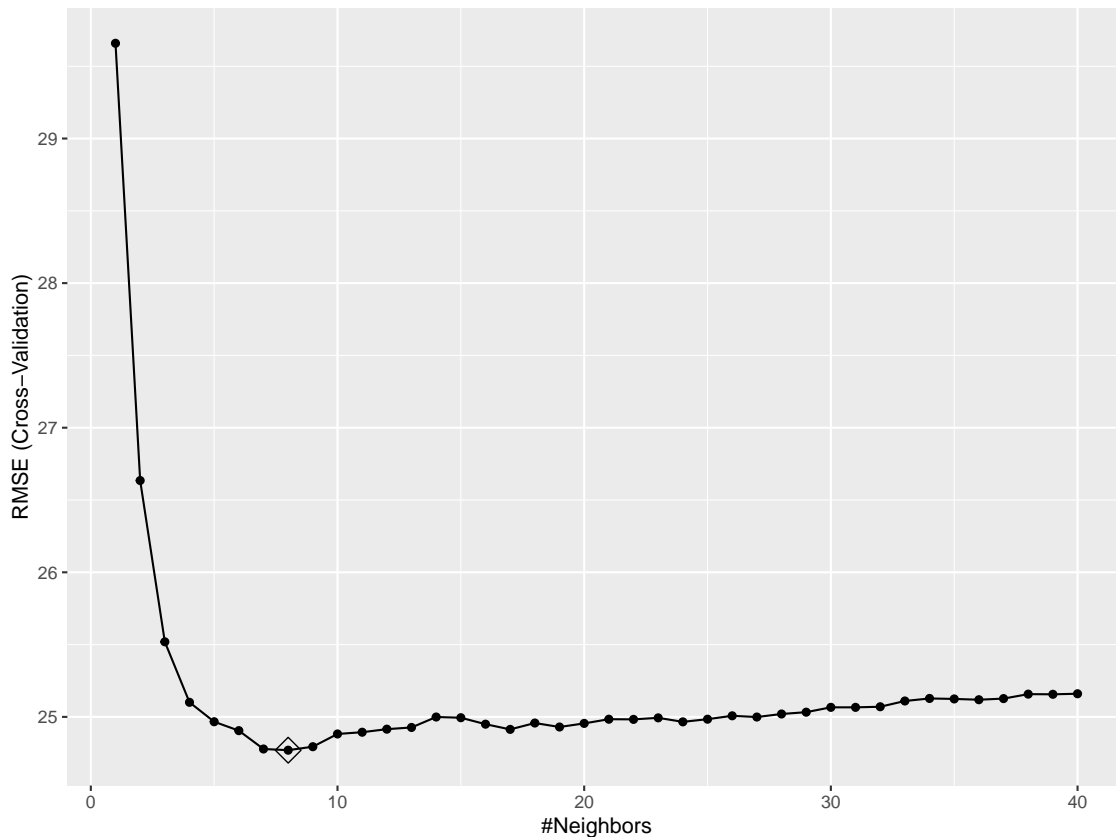
```
## Residuals:
##      Min       1Q   Median       3Q      Max
## -73.848 -13.472  -1.414   9.399 257.759
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -2.743e+03  1.295e+02 -21.171 < 2e-16 ***
## age         -1.353e-02  1.177e-01  -0.115 0.908505
## gender1     -4.538e+00  9.432e-01  -4.812 1.59e-06 ***
## race2        5.460e-01  2.166e+00   0.252 0.800997
## race3       -1.639e+00  1.193e+00  -1.374 0.169550
## race4       -9.671e-01  1.653e+00  -0.585 0.558518
## smoking1     3.052e+00  1.066e+00   2.863 0.004236 **
## smoking2     7.843e+00  1.563e+00   5.018 5.58e-07 ***
## height       1.608e+01  7.612e-01  21.120 < 2e-16 ***
## weight      -1.733e+01  8.043e-01 -21.541 < 2e-16 ***
## bmi          5.195e+01  2.300e+00  22.587 < 2e-16 ***
## hypertension1 3.229e+00  1.560e+00   2.070 0.038547 *
## diabetes1    -2.210e-01  1.277e+00  -0.173 0.862608
## SBP          3.232e-04  1.012e-01   0.003 0.997453
## LDL         -1.013e-02  2.499e-02  -0.405 0.685287
## vaccine1     -8.167e+00  9.684e-01  -8.434 < 2e-16 ***
## severity1     5.717e+00  1.534e+00   3.726 0.000199 ***
## studyB        4.315e+00  1.215e+00   3.552 0.000389 ***
## studyC       -3.479e-01  1.500e+00  -0.232 0.816566
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 23.58 on 2498 degrees of freedom
## Multiple R-squared:  0.2754, Adjusted R-squared:  0.2702
## F-statistic: 52.75 on 18 and 2498 DF, p-value: < 2.2e-16

# view performance on the test set (RMSE)
lm_test_pred <- predict(linear_model, newdata = test_x) # test dataset
lm_test_rmse <- sqrt(mean((lm_test_pred - test_y)^2))
sprintf("test error for lm is: %.3f",lm_test_rmse)

## [1] "test error for lm is: 23.569"
```

## K-Nearest Neighbors (KNN)

```
set.seed(8)
# fit knn on train data use caret
kGrid <- expand.grid(k = seq(1, to = 40, by = 1))
knn_model <- train(train_x,
                   train_y,
                   method = "knn",
                   trControl = ctrl,
                   tuneGrid = kGrid)
ggplot(knn_model, highlight = TRUE)
```



```
# knn with K = 18 was selected as the final model
```

```
# view performance on the test set (RMSE)
```

```
knn_test_pred <- predict(knn_model, newdata = test_x) # test dataset
knn_test_rmse <- sqrt(mean((knn_test_pred - test_y)^2))
sprintf("test error for K-Nearest Neighbors is: %.3f", knn_test_rmse)
```

```
## [1] "test error for K-Nearest Neighbors is: 25.181"
```

## Enet

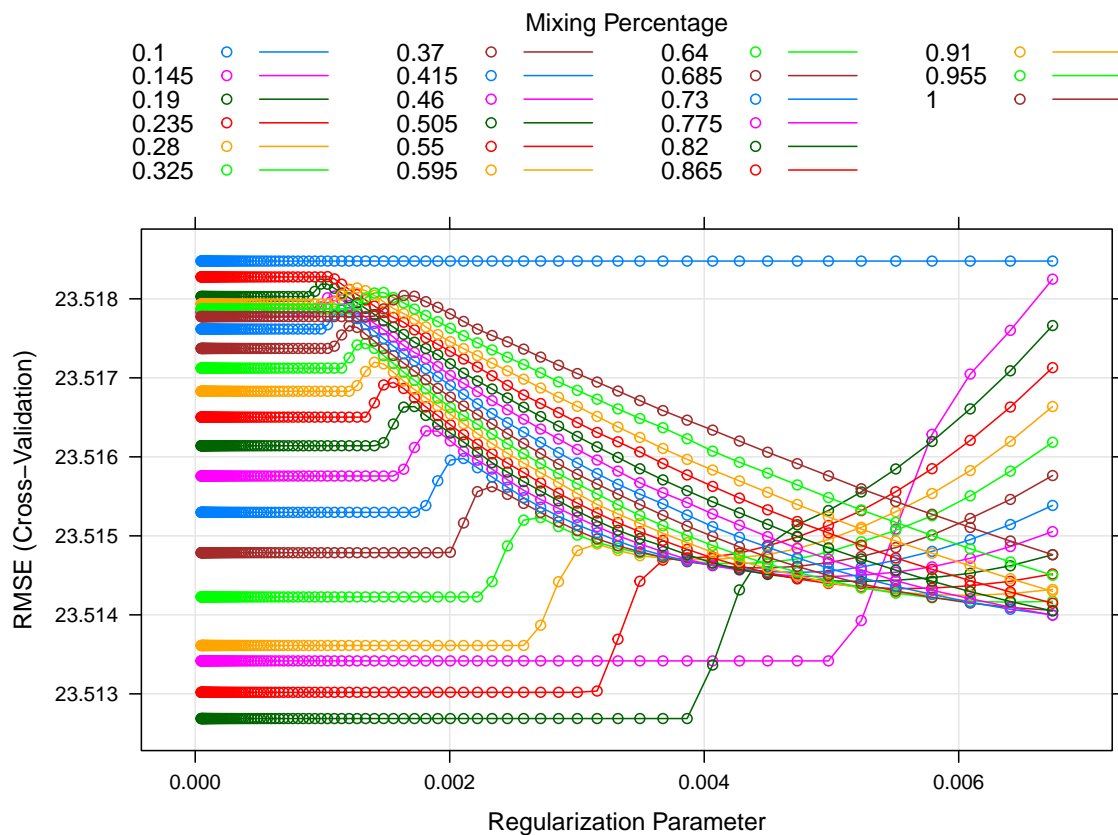
```
#tuning glmnet alpha 0 and 0.05 have large RMSE. So alpha start from 0.1
```

```
set.seed(8)
```

```
enet.fit.min <- train(train_x, train_y,
                      method = "glmnet",
                      tuneGrid = expand.grid(alpha = seq(0.1, 1, length = 21),
                                             lambda = exp(seq(-10, -5, length = 100))),
                      trControl = ctrl)
enet.fit.min$bestTune
```

```
##      alpha      lambda
## 289  0.19 0.00386592
```

```
plot(enet.fit.min)
```

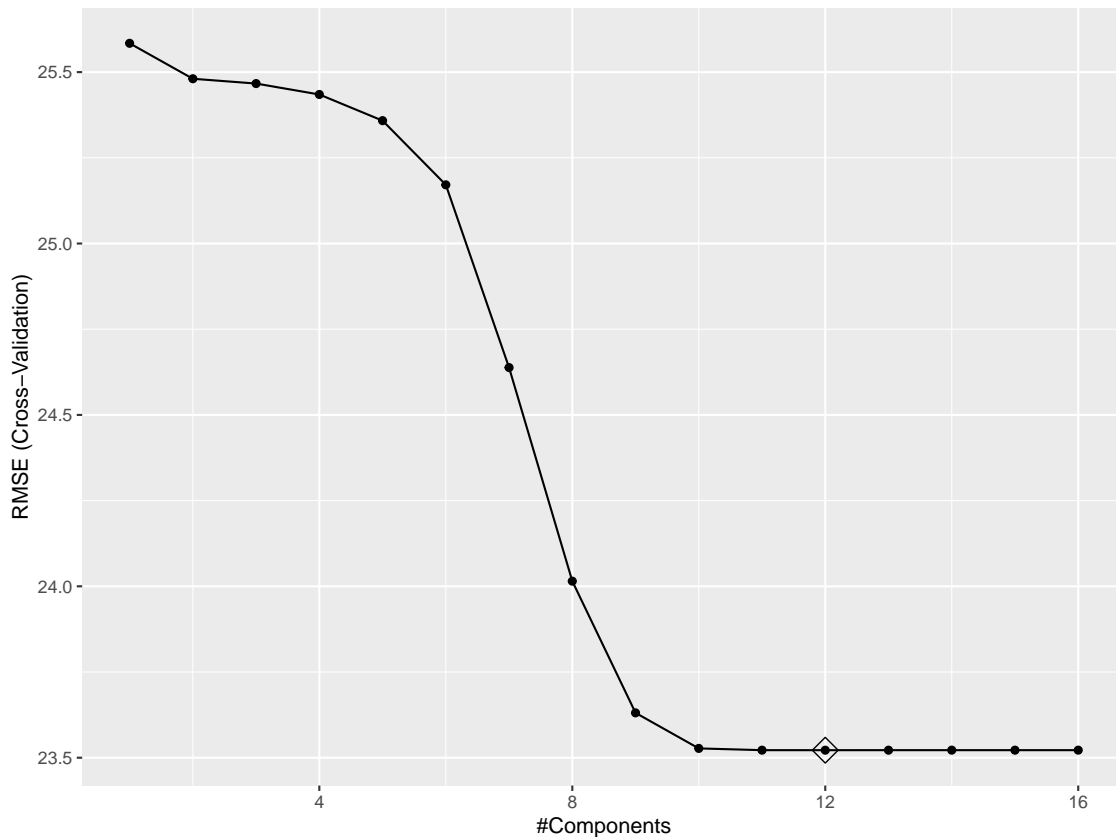


```
# view performance on the test set (RMSE)
enet_test_pred <- predict(enet.fit.min, newdata = test_x) # test dataset
enet_test_rmse <- sqrt(mean((enet_test_pred - test_y)^2))
sprintf("test error for enet is: %.3f", enet_test_rmse)
```

```
## [1] "test error for enet is: 23.523"
```

## PLS

```
set.seed(8)
pls_model <- train(train_x,
  train_y,
  method = "pls",
  tuneGrid = data.frame(ncomp = 1:16),
  trControl = ctrl,
  preProcess = c("center", "scale"))
ggplot(pls_model, highlight = TRUE)
```



```
ggsave(file = "image/pls_number_of_component.png", width = 10, height = 7)
```

```
# view performance on the test set (RMSE)
pls_test_pred <- predict(pls_model, newdata = test_x) # test dataset
pls_test_rmse <- sqrt(mean((pls_test_pred - test_y)^2))
sprintf("test error for pls is: %.3f", pls_test_rmse)
```

```
## [1] "test error for pls is: 23.569"
```

## Generalised additive regression (GAM)

```
# gam use caret
# fit GAM model using all predictors
# parallel computing
no_cores <- detectCores() - 1
cl <- makePSOCKcluster(no_cores)
registerDoParallel(cl)
set.seed(8)
gam_model <- train(train_x, train_y, # training dataset
                   method = "gam",
                   trControl = ctrl,
                   control = gam.control(maxit = 200))

stopCluster(cl)
registerDoSEQ()

gam_model$bestTune
```

```
## select method
```

```
## 2 TRUE GCV.Cp
```

```
gam_model$finalModel
```

```
##
```

```
## Family: gaussian
```

```
## Link function: identity
```

```
##
```

```
## Formula:
```

```
## .outcome ~ gender1 + race2 + race3 + race4 + smoking1 + smoking2 +
```

```
## hypertension1 + diabetes1 + vaccine1 + severity1 + studyB +
```

```
## studyC + s(age) + s(SBP) + s(LDL) + s(bmi) + s(height) +
```

```
## s(weight)
```

```
##
```


```
## Estimated degrees of freedom:
```

```
## 0.00 0.00 0.00 7.78 2.65 4.50 total = 27.93
```

```
##
```

```
## GCV score: 448.2538
```

```
plot(gam_model)
```



```
summary(gam_model$finalModel)
```

```
##
```

```
## Family: gaussian
```

```
## Link function: identity
```

```
##
```

```
## Formula:
```

```
## .outcome ~ gender1 + race2 + race3 + race4 + smoking1 + smoking2 +
```

```
## hypertension1 + diabetes1 + vaccine1 + severity1 + studyB +
```

```
##      studyC + s(age) + s(SBP) + s(LDL) + s(bmi) + s(height) +
##      s(weight)
##
## Parametric coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   43.7183     1.3232  33.040 < 2e-16 ***
## gender1       -4.6981     0.8434  -5.570 2.82e-08 ***
## race2          0.5405     1.9332   0.280 0.779829
## race3         -1.5605     1.0659  -1.464 0.143332
## race4         -2.6252     1.4791  -1.775 0.076054 .
## smoking1       4.3355     0.9556   4.537 5.98e-06 ***
## smoking2       8.0325     1.3975   5.748 1.01e-08 ***
## hypertension1  3.0143     0.8471   3.558 0.000380 ***
## diabetes1      0.6857     1.1425   0.600 0.548453
## vaccine1      -7.7120     0.8649  -8.916 < 2e-16 ***
## severity1      5.9722     1.3716   4.354 1.39e-05 ***
## studyB         4.1736     1.0846   3.848 0.000122 ***
## studyC        -0.6406     1.3400  -0.478 0.632678
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Approximate significance of smooth terms:
##              edf Ref.df       F p-value
## s(age)        1.123e-06    9   0.000 0.55847
## s(SBP)        8.397e-07    9   0.000 0.68712
## s(LDL)        1.176e-06    9   0.000 0.51661
## s(bmi)        7.775e+00    9 111.700 < 2e-16 ***
## s(height)     2.649e+00    9   1.019 0.00269 **
## s(weight)     4.501e+00    9   1.049 0.04744 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## R-sq.(adj) =  0.418   Deviance explained = 42.5%
## GCV = 448.25   Scale est. = 443.28    n = 2517
```

```
# view performance on the test set (RMSE)
gam_test_pred <- predict(gam_model, newdata = test_x) # test dataset
gam_test_rmse <- sqrt(mean((gam_test_pred - test_y)^2))
sprintf("test error for Generalised additive regression is: %.3f", knn_test_rmse)
```

```
## [1] "test error for Generalised additive regression is: 25.181"
```

## Multivariate adaptive regression

```
set.seed(8)

# create dummy variables for categorical variables
df_dummies <- data.frame(model.matrix(~ . - 1,
                                     # exclude ID and continuous variables
                                     data = data[, c("gender", "race", "smoking", "hypertension", "diabetes",
                                                       "vaccine", "severity", "study")]),
                          # add continuous variables back to the data frame
                          age = data$age,
                          height = data$height,
```

```

weight = data$weight,
bmi = data$bmi,
SBP = data$SBP,
LDL = data$LDL,
recovery_time = data$recovery_time)

# rename df_dummies dataset as dat
dat_mars <- df_dummies

# training data, using previous indexTrain so the same as x_train
dat_train_mars <- dat_mars[indexTrain, ]
x_mars <- model.matrix(recovery_time~.,dat_mars)[indexTrain,-1]
y_mars <- dat_mars$recovery_time[indexTrain]

# test data, using previous indexTrain so the same as x_test
dat_test_mars <- dat_mars[-indexTrain, ]
x2_mars <- model.matrix(recovery_time~.,dat_mars)[-indexTrain,-1]
y2_mars <- dat_mars$recovery_time[-indexTrain]

# create grid of all possible pairs that can take degree and nprune values
mars_grid <- expand.grid(degree = 1:3, # number of possible product hinge functions in 1 term
                        nprune = 2:18) # upper bound of number of terms in model


# parallel computing
cl <- makePSOCKcluster(no_cores)
registerDoParallel(cl)

set.seed(8)
mars_model <- train(x_mars, y_mars, # training dataset
                    method = "earth",
                    tuneGrid = mars_grid,
                    trControl = ctrl)

stopCluster(cl)
registerDoSEQ()


ggplot(mars_model)

```



```
print(plot(mars_model))
```





```
summary(mars_model$finalModel)
```

```
## Call: earth(x=matrix[2517,19], y=c(14,36,50,65,3...), keepxy=TRUE, degree=2,
##          nprune=5)
```

```
##
##          coefficients
## (Intercept)      -1.264634
## vaccine1         -7.897168
## h(bmi-25.9)       8.611315
## h(32.6-bmi)       5.699561
## studyB * h(bmi-32.6) 41.724058
##
```

```
## Selected 5 of 24 terms, and 3 of 19 predictors (nprune=5)
```

```
## Termination condition: Reached nk 39
```

```
## Importance: studyB, bmi, vaccine1, gender0-unused, gender1-unused, ...
```

```
## Number of terms at each degree of interaction: 1 3 1
```

```
## GCV 413.5903    RSS 1031928    GRSq 0.4574914    RSq 0.4617953
```

```
# view performance on the test set (RMSE)
```

```
mars_test_pred <- predict(mars_model, newdata = x2_mars) # test dataset
```

```
mars_test_rmse <- sqrt(mean((mars_test_pred - test_y)^2))
```

```
sprintf("test error for MARS is: %.3f", mars_test_rmse)
```

```
## [1] "test error for MARS is: 22.586"
```

## Random Forest - Regression

```

set.seed(8)
rf.grid <- expand.grid(mtry = 1:14,
                      splitrule = "variance",
                      min.node.size = seq(from = 1, to = 50, by = 5))

# parallel computing
no_cores <- detectCores() - 1
cl <- makePSOCKcluster(no_cores)
registerDoParallel(cl)

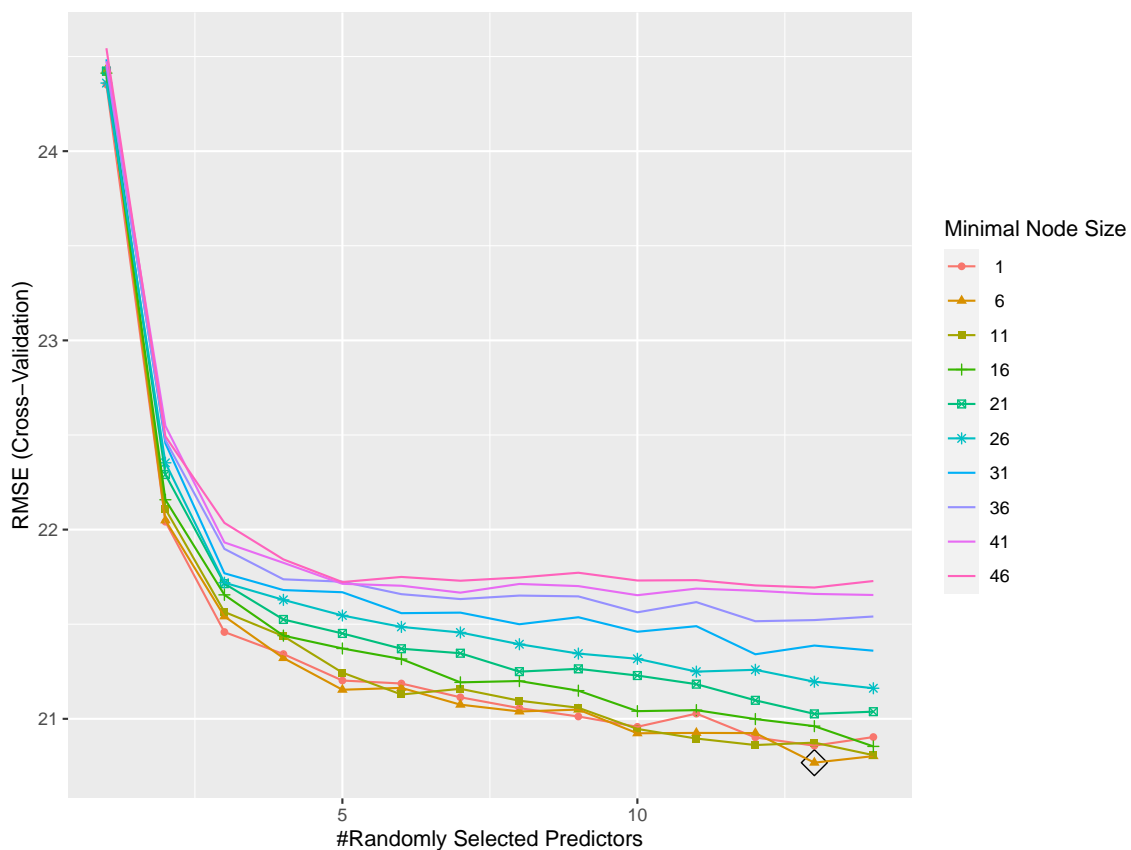
set.seed(8)
rf_model <- train(recovery_time ~ . ,
                  train_data,
                  method = "ranger",
                  tuneGrid = rf.grid,
                  trControl = ctrl)

stopCluster(cl)
registerDoSEQ()
ggplot(rf_model, highlight = TRUE)

```

## Warning: The shape palette can deal with a maximum of 6 discrete values because  
 ## more than 6 becomes difficult to discriminate; you have 10. Consider  
 ## specifying shapes manually if you must have them.

## Warning: Removed 56 rows containing missing values (`geom\_point()`).



```
rf_model$bestTune
```

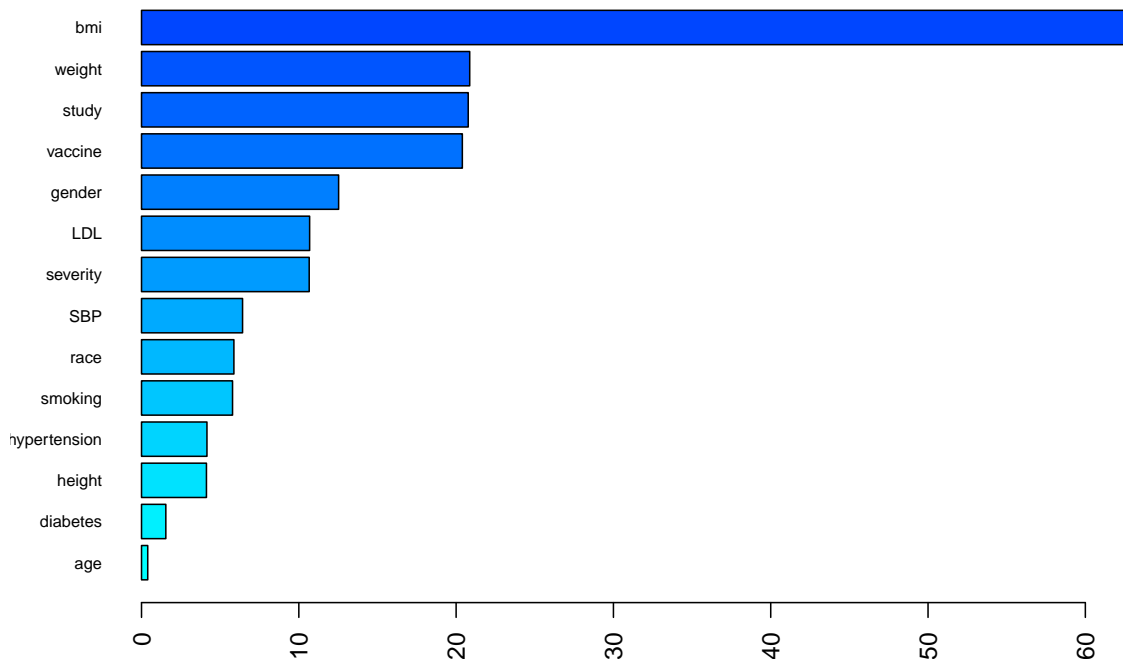
```
##      mtry splitrule min.node.size  
## 122    13  variance              6
```

```
# variable importance
```

```
set.seed(8)
```

```
rf.final.per <- ranger(recovery_time ~ . ,  
                      train_data,  
                      mtry = rf_model$bestTune[[1]],  
                      splitrule = "variance",  
                      min.node.size = rf_model$bestTune[[3]],  
                      importance = "permutation",  
                      scale.permutation.importance = TRUE)
```

```
barplot(sort(ranger::importance(rf.final.per), decreasing = FALSE),  
        las = 2, horiz = TRUE, cex.names = 0.7,  
        col = colorRampPalette(colors = c("cyan", "blue"))(19))
```



```
# view performance on the test set (RMSE)
```

```
rf_test_pred <- predict(rf_model, newdata = test_data) # test dataset
```

```
rf_test_rmse <- sqrt(mean((rf_test_pred - test_y)^2))
```

```
sprintf("test error for Random Forest is: %.3f", rf_test_rmse)
```

```
## [1] "test error for Random Forest is: 21.965"
```

## Boosting

```
gbm.grid <- expand.grid(n.trees = c(500,1000,2000,3000,4000,5000),  
                      interaction.depth = 1:5,  
                      shrinkage = c(0.005,0.01),  
                      n.minobsinnode = c(1))
```

```
no_cores <- detectCores() - 1
```

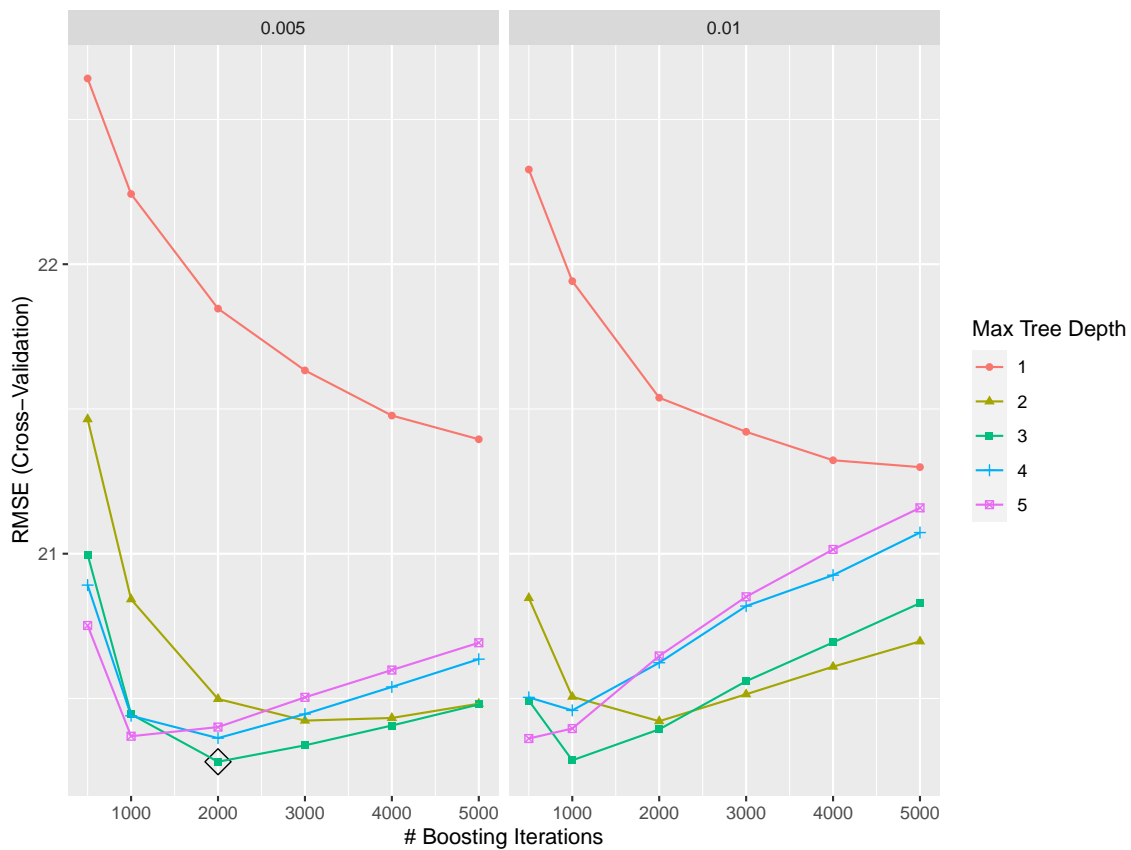
```

c1 <- makePSOCKcluster(no_cores)
registerDoParallel(c1)
set.seed(8)
gbm.fit = train(recovery_time~. ,
                train_data,
                tuneGrid = gbm.grid,
                trControl = ctrl,
                method = "gbm",
                verbose = FALSE)

stopCluster(c1)
registerDoSEQ()

ggplot(gbm.fit, highlight = TRUE)

```



```
gbm.fit$bestTune
```

```
##      n.trees interaction.depth shrinkage n.minobsinnode
## 15      2000                3      0.005                1
```

```
# train error
```

```
gbm.fit.predict = predict(gbm.fit, newdata = train_data)
RMSE(gbm.fit.predict, train_y)
```

```
## [1] 17.48103
```

```
# test error
```

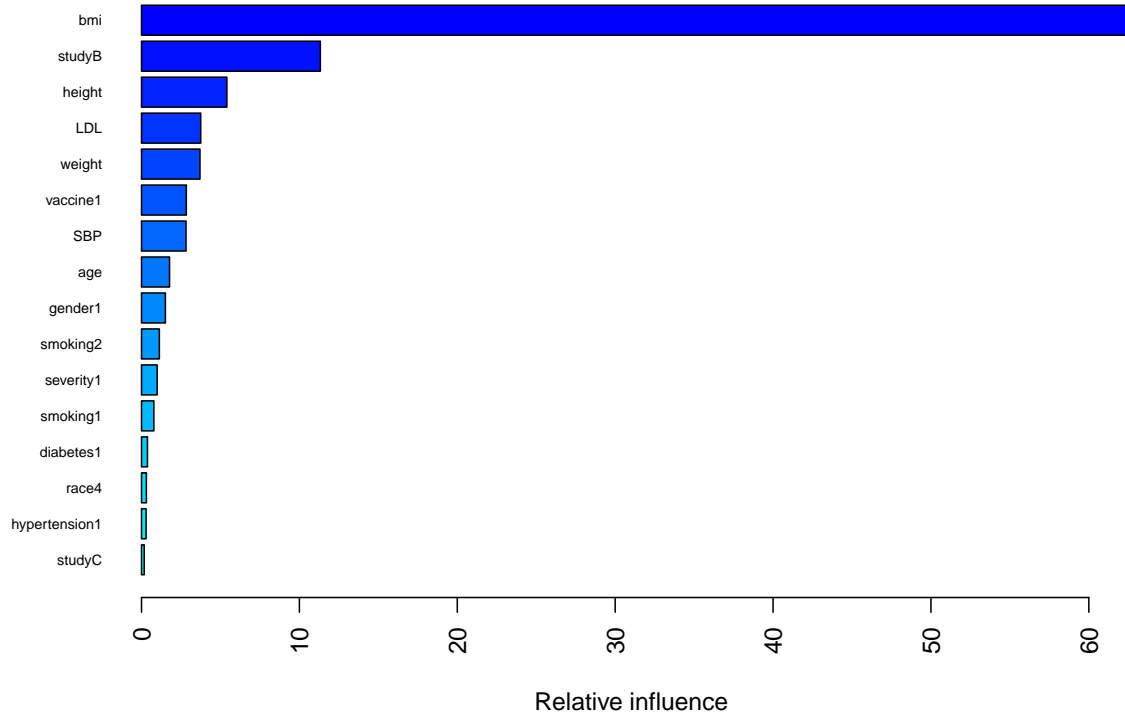
```
gbm.fit.predict = predict(gbm.fit, newdata = test_data)
```

```
RMSE(gbm.fit.predict, test_y)
```

```
## [1] 21.36644
```

```
#vip
```

```
summary(gbm.fit$finalModel, las = 2, cBars = 16, cex.names = 0.6)
```



```
##           var      rel.inf
## bmi          bmi 62.67390364
## studyB        studyB 11.32550304
## height        height  5.40662442
## LDL           LDL   3.74887726
## weight        weight  3.70130262
## vaccine1      vaccine1  2.83797028
## SBP           SBP   2.81893398
## age           age   1.77559187
## gender1       gender1  1.51170747
## smoking2      smoking2  1.12956594
## severity1     severity1  0.98910314
## smoking1      smoking1  0.78002955
## diabetes1     diabetes1  0.37707597
## race4         race4   0.30966831
## hypertension1 hypertension1 0.28792897
## studyC        studyC   0.17053920
## race3         race3   0.09956841
## race2         race2   0.05610592
```

```
#pdp
```

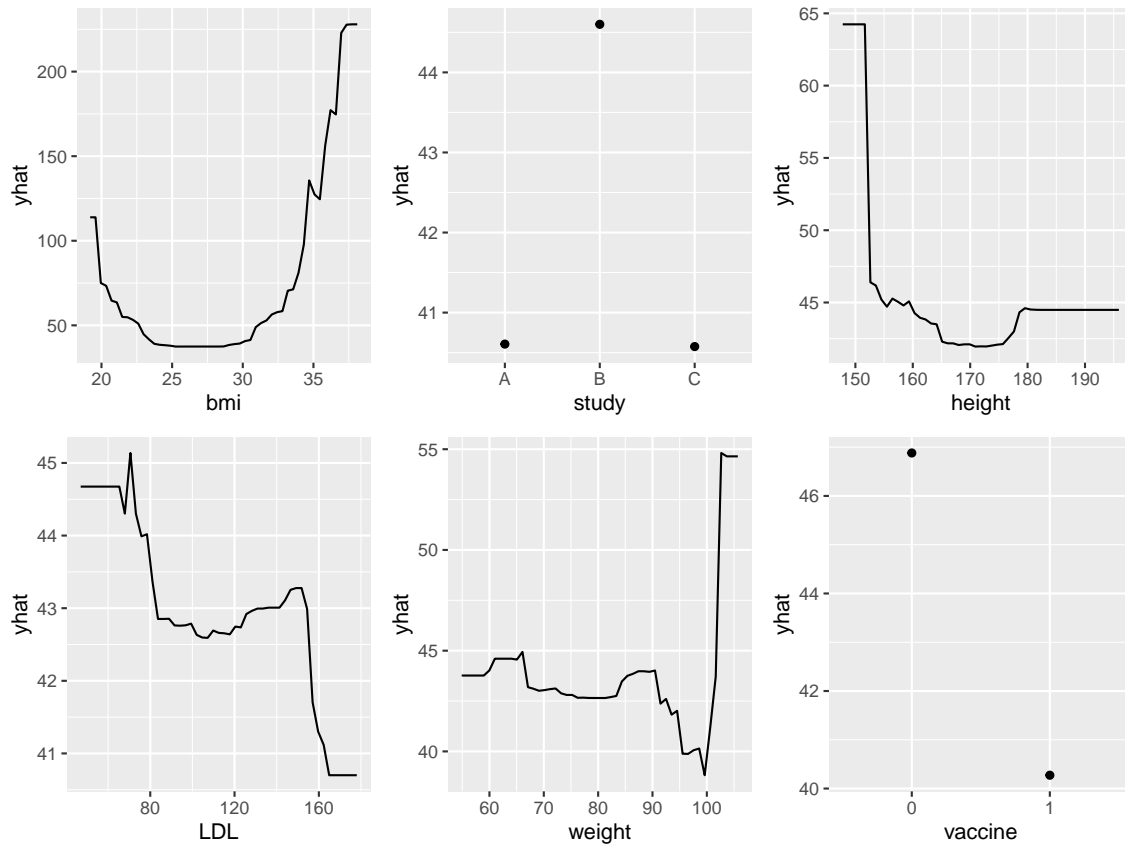
```
features = c("bmi", "study", "height", "LDL", "weight", "vaccine")
```

```
pdp <- lapply(features, function(x) {
  partial(gbm.fit, pred.var = x, which.class = 2,
```

```

    prob = TRUE, plot = TRUE, plot.engine = "ggplot2")
  })
  grid.arrange(grobs = pdps, ncol = 3)

```



## resampling on train

```

# parallel computing
no_cores <- detectCores() - 1
cl <- makePSOCKcluster(no_cores)
registerDoParallel(cl)
set.seed(8)

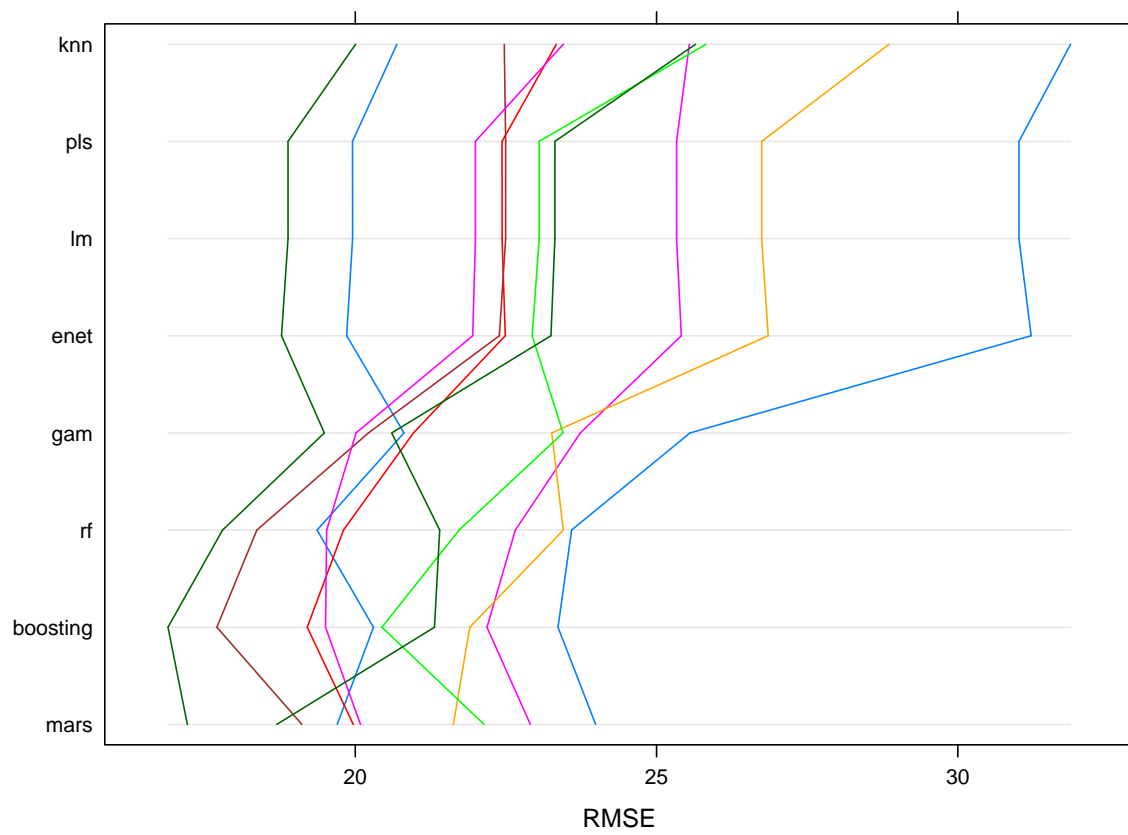
resamp <- resamples(list(lm = linear_model,
                        knn = knn_model,
                        enet = enet.fit.min,
                        pls = pls_model,
                        gam = gam_model,
                        mars = mars_model,
                        rf = rf_model,
                        boosting = gbm.fit))

stopCluster(cl)
registerDoSEQ()

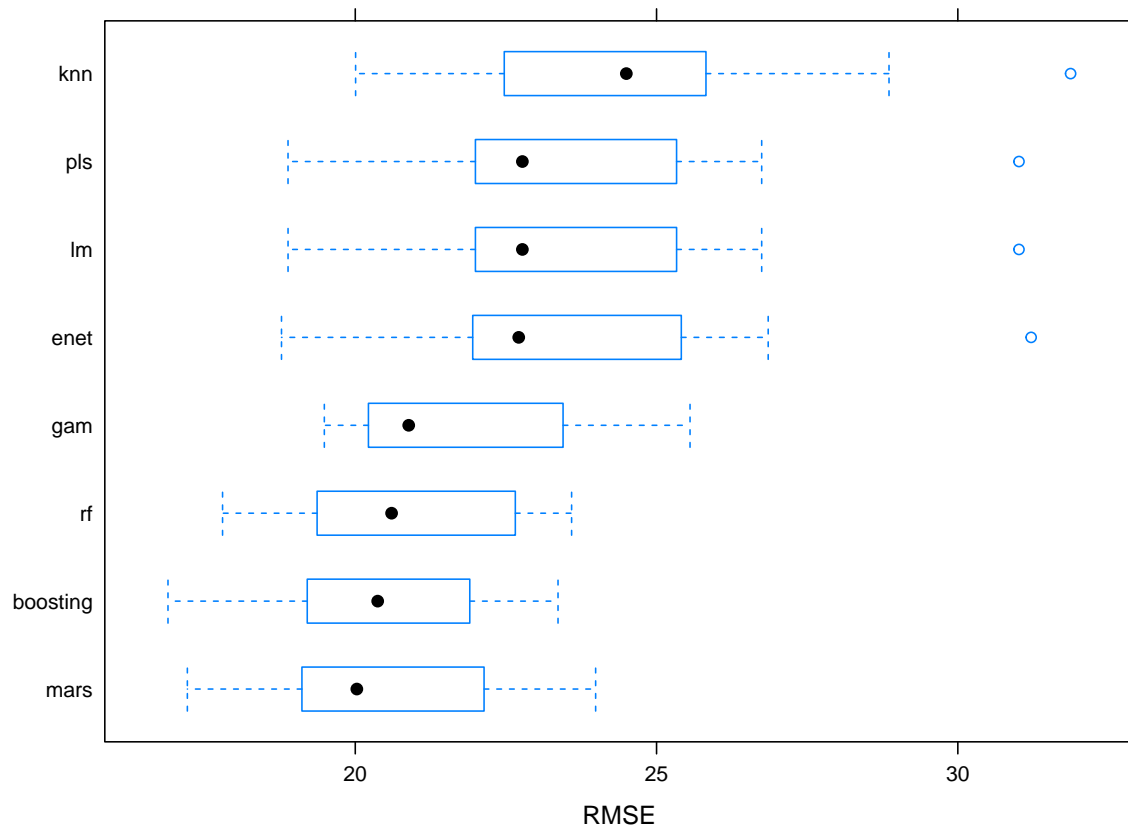
summary(resamp)

```

```
##
## Call:
## summary.resamples(object = resamp)
##
## Models: lm, knn, enet, pls, gam, mars, rf, boosting
## Number of resamples: 10
##
## MAE
##           Min.   1st Qu.   Median     Mean   3rd Qu.     Max. NA's
## lm       13.76934 15.43334 15.76924 15.80693 16.17432 18.11764    0
## knn       14.06701 15.73153 16.09257 16.03670 16.46648 17.89633    0
## enet      13.69102 15.32867 15.68785 15.72506 16.12487 18.10713    0
## pls       13.76888 15.43393 15.76897 15.80691 16.17445 18.11772    0
## gam       13.19262 14.38982 14.82480 14.68947 15.13537 15.73301    0
## mars      12.27144 14.08073 14.54892 14.38932 14.86938 15.47013    0
## rf        13.28202 13.87283 14.49171 14.30696 14.74181 15.15497    0
## boosting  12.32434 13.76258 14.07635 13.93779 14.36728 15.05691    0
##
## RMSE
##           Min.   1st Qu.   Median     Mean   3rd Qu.     Max. NA's
## lm       18.88430 22.10438 22.77336 23.52201 24.82757 31.01445    0
## knn       20.00557 22.68863 24.49846 24.76979 25.77652 31.87127    0
## enet      18.77573 22.05979 22.71284 23.51269 24.86986 31.21659    0
## pls       18.88388 22.10432 22.77337 23.52193 24.82749 31.01461    0
## gam       19.48652 20.31487 20.88715 21.80966 23.40153 25.55548    0
## mars      17.21334 19.26111 20.02599 20.54374 22.01000 23.98902    0
## rf        17.79759 19.40633 20.60269 20.76831 22.42284 23.59035    0
## boosting  16.89195 19.27883 20.37222 20.28145 21.75394 23.36495    0
##
## Rsquared
##           Min.   1st Qu.   Median     Mean   3rd Qu.     Max. NA's
## lm       0.15262099 0.2417259 0.2664909 0.2612246 0.2814750 0.3695876    0
## knn       0.06890501 0.1229402 0.1601685 0.1838074 0.1914104 0.3999299    0
## enet      0.15157916 0.2419275 0.2660519 0.2606718 0.2821525 0.3650478    0
## pls       0.15266199 0.2417286 0.2664926 0.2612288 0.2814784 0.3695807    0
## gam       0.19032422 0.3458940 0.3645476 0.3822057 0.4249498 0.6033048    0
## mars      0.17813209 0.3790794 0.4295123 0.4310183 0.5145510 0.5980581    0
## rf        0.23449216 0.3940702 0.4288848 0.4227968 0.4520670 0.6156429    0
## boosting  0.21179764 0.4242251 0.4766269 0.4578489 0.5023122 0.6160659    0
parallelplot(resamp, metric = "RMSE")
```



```
bwplot(resamp, metric = "RMSE")
```

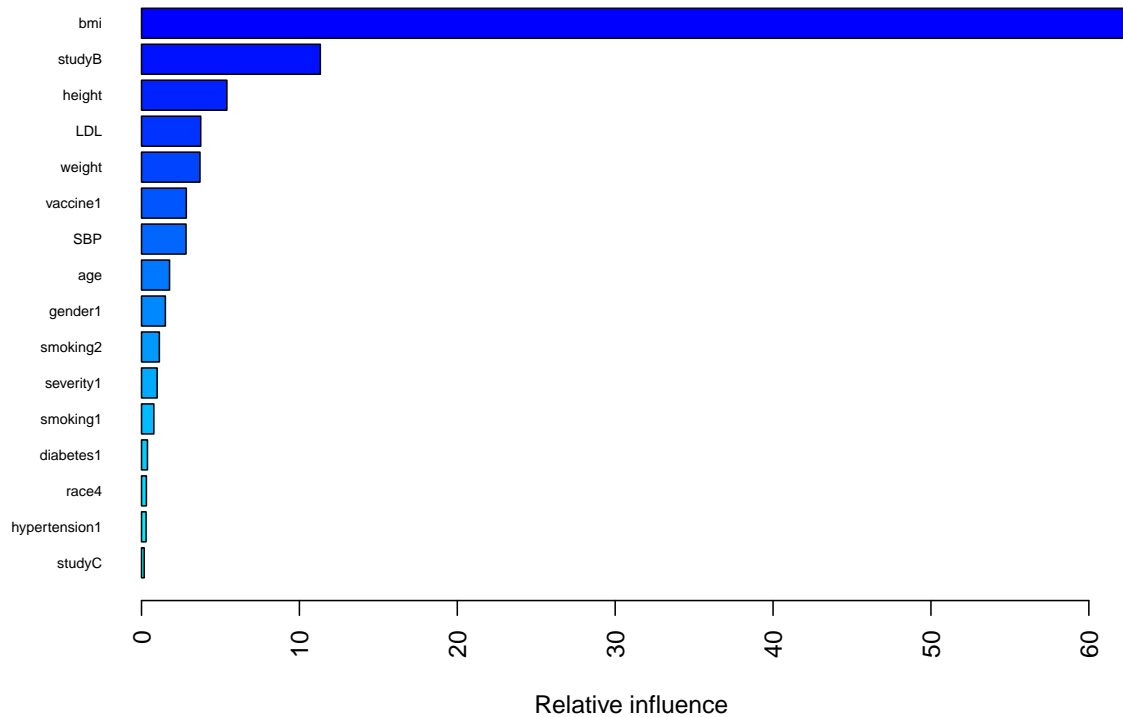




## Final Model for Regression

```
summary(gbm.fit$finalModel, las = 2, cBars = 16, cex.names = 0.6)
```

### PDP plot



```
##           var      rel.inf
## bmi          bmi 62.67390364
## studyB       studyB 11.32550304
## height      height  5.40662442
## LDL          LDL  3.74887726
## weight      weight  3.70130262
## vaccine1    vaccine1  2.83797028
## SBP          SBP  2.81893398
## age          age  1.77559187
## gender1     gender1  1.51170747
## smoking2    smoking2  1.12956594
## severity1   severity1  0.98910314
## smoking1    smoking1  0.78002955
## diabetes1   diabetes1  0.37707597
## race4       race4  0.30966831
## hypertension1 hypertension1 0.28792897
## studyC      studyC  0.17053920
## race3       race3  0.09956841
## race2       race2  0.05610592
```


```
#pdp
features = c("bmi", "study", "height", "LDL", "weight", "vaccine")

pdps <- lapply(features, function(x) {
  partial(gbm.fit, pred.var = x, which.class = 2,
```

```

    prob = TRUE, plot = TRUE, plot.engine = "ggplot2")
  })
  grid.arrange(grobs = pdps, ncol = 3)

```

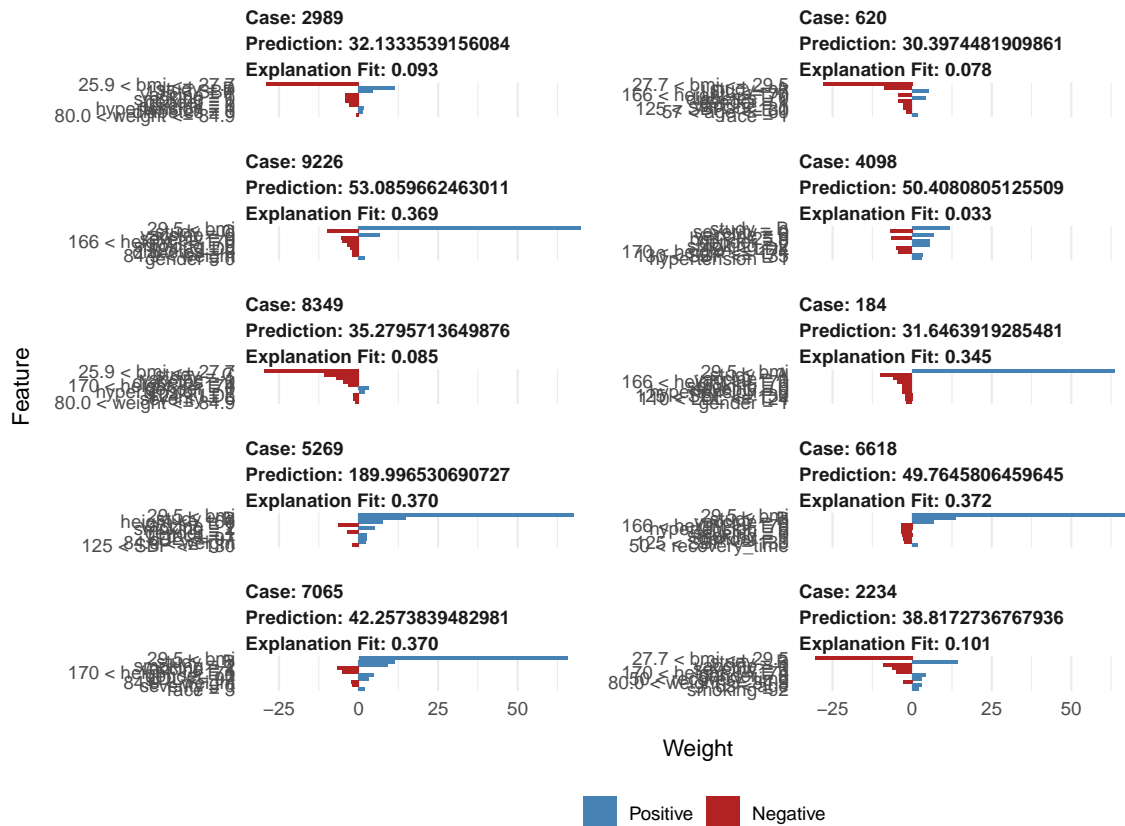


```

#lime
explain_boosting <- lime(train_data, gbm.fit)
new_obs <- explain(train_data[1:10,], explain_boosting, n_features=10)

plot_features(new_obs)

```



## Classification

### Data Manipulation

We consider time to recovery as a binary outcome as long(>30 days) and short (<= 30 days).

```
#consider time to recovery as a binary outcome (>30 days vs. <= 30 days) and develop a prediction model
train_class =
  train_data %>%
  mutate(recovery_time = ifelse(recovery_time>30, "long", "short"))
test_class =
  test_data %>%
  mutate(recovery_time = ifelse(recovery_time>30, "long", "short"))

# matrix of predictors
x_train_class <- model.matrix(recovery_time~.,train_class)[,-1]
# vector of response
y_train_class <- train_class$recovery_time

# matrix of predictors
x_test_class <- model.matrix(recovery_time~.,test_class)[,-1]
# vector of response
y_test_class <- test_class$recovery_time

# ctrl for class
ctrl_class <- trainControl(method = "cv",
```


```
summaryFunction = twoClassSummary,  
classProbs = TRUE)
```

## Modeling Strategy

- Penalized logistic regression
- Generalized additive model (GAM)
- Multivariate adaptive regression splines (MARS)
- Linear discriminant analysis (LDA)
- Random Forest
- Boosting
- Support Vector Machine (SVM with Linear and Radial Kernels)

## Penalized logistic regression

```
glmnetGrid <- expand.grid(.alpha = seq(0, 1, length = 21),  
                          .lambda = exp(seq(-12, -7, length = 50)))  
set.seed(8)  
model.glmnet <- train(x = x_train_class,  
                      y = y_train_class,  
                      method = "glmnet",  
                      tuneGrid = glmnetGrid,  
                      metric = "ROC",  
                      trControl = ctrl_class)  
  
model.glmnet$bestTune  
  
##      alpha      lambda  
## 524    0.5 6.422865e-05  
  
myCol <- rainbow(25)  
myPar <- list(superpose.symbol = list(col = myCol),  
              superpose.line = list(col = myCol))  
  
plot(model.glmnet, par.settings = myPar, xTrans = function(x) log(x))
```



```
# view performance on the test set (accuracy)
glmn_class_test_pred <- predict(model.glmn, newdata=x_test_class) # test dataset
glmn_class_test_acc <- sum(glmn_class_test_pred==y_test_class)/length(y_test_class)
sprintf("test error for penalized logistic is: %.3f", 1-glmn_class_test_acc)
```

```
## [1] "test error for penalized logistic is: 0.286"
```

Since the boundary value for  $\lambda$  is  $\exp(-10)=4.54e-5$  and  $\exp(-5)=0.007$ , and the optimal lambda is not on the boundary. So local minimum is achieved.

## Gam

```
# gam
set.seed(8)
model.gam <- train(x = x_train_class,
                  y = y_train_class,
                  method = "gam",
                  tuneGrid = data.frame(method = "GCV.Cp", select = c(TRUE,FALSE)),
                  metric = "ROC",
                  trControl = ctrl_class)
```

```
model.gam$bestTune
```

```
## select method
```

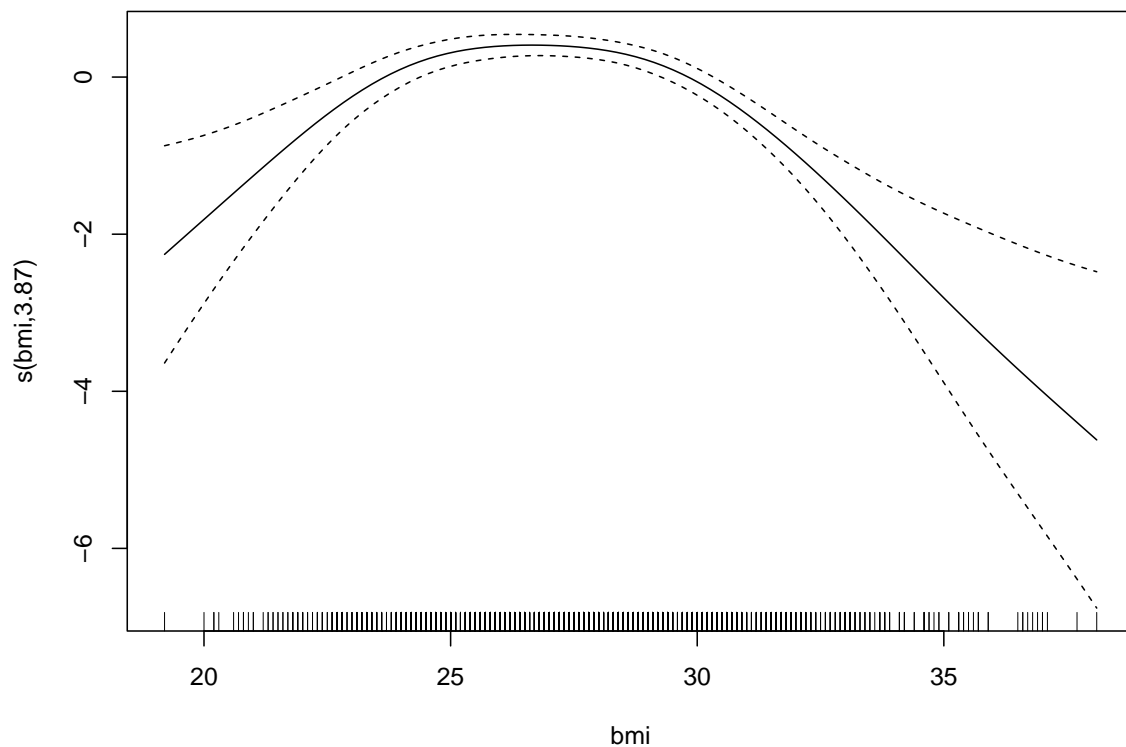
```
## 2 TRUE GCV.Cp
```

```
model.gam$finalModel
```

```
##
```

```
## Family: binomial
## Link function: logit
##
## Formula:
## .outcome ~ gender1 + race2 + race3 + race4 + smoking1 + smoking2 +
##     hypertension1 + diabetes1 + vaccine1 + severity1 + studyB +
##     studyC + s(age) + s(SBP) + s(LDL) + s(bmi) + s(height) +
##     s(weight)
##
## Estimated degrees of freedom:
## 0.0002 3.0898 0.1982 3.8717 1.9808 0.0001 total = 22.14
##
## UBRE score: 0.07604226
```

```
#discover bmi and recovery time
plot(model.gam$finalModel, select = 4)
```



```
#train error
train.pred.gam <- predict(model.gam, newdata = x_train_class)
train.error.gam = 1-sum(train_class$recovery_time == train.pred.gam)/length(train_data$recovery_time)
sprintf("The train error for GAM is %.3f", train.error.gam)
```

```
## [1] "The train error for GAM is 0.266"
```

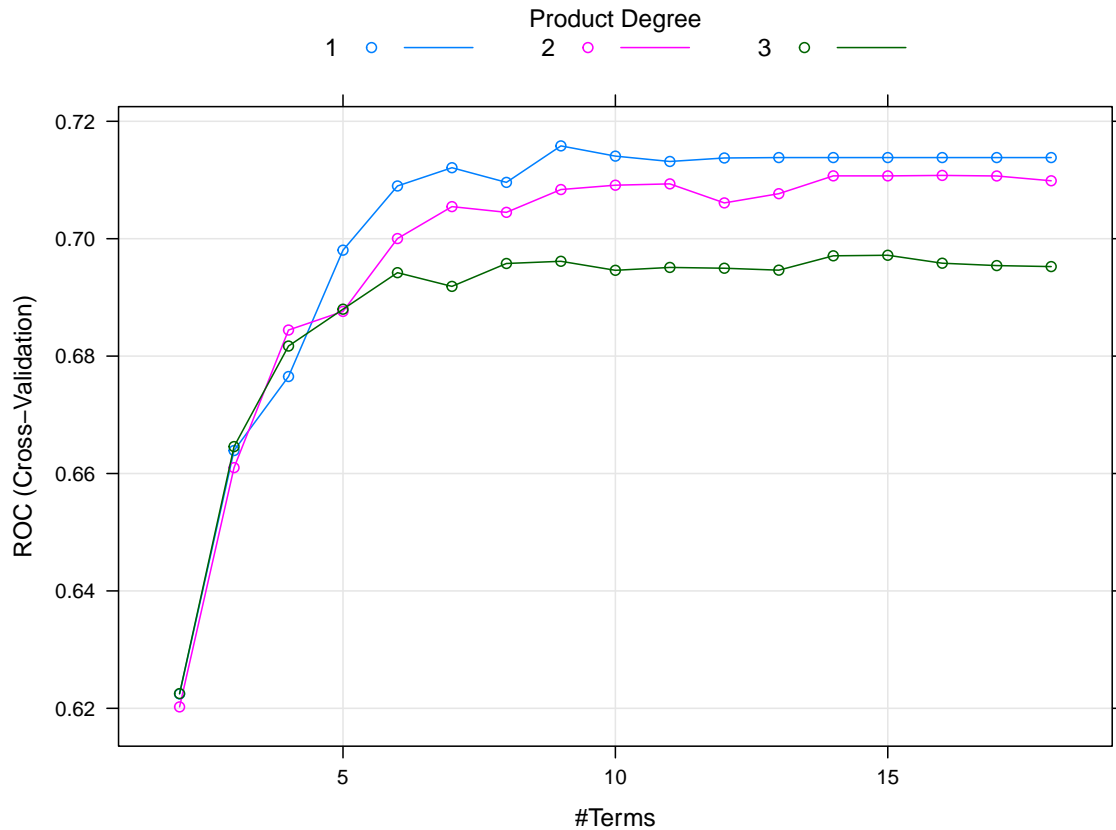
```
#test error
test.pred.gam <- predict(model.gam, newdata = x_test_class)
test.error.gam = 1-sum(test_class$recovery_time == test.pred.gam)/length(test_class$recovery_time)
sprintf("The test error for GAM is %.3f", test.error.gam)
```

```
## [1] "The test error for GAM is 0.283"
```

## MARS

```
set.seed(8)
mars.fit = train(x = x_train_class,
                 y = y_train_class,
                 method = "earth",
                 tuneGrid = expand.grid(degree = 1:3,
                                       nprune = 2:18),
                 metric = "ROC",
                 trControl = ctrl_class)

plot(mars.fit)
```



```
#train error
train.pred.mars = predict(mars.fit , newdata = x_train_class)
train.error.mars = 1-sum(train_class$recovery_time == train.pred.mars)/length(train_data$recovery_time)
sprintf("The train error for MARS is %.3f", train.error.mars)
```

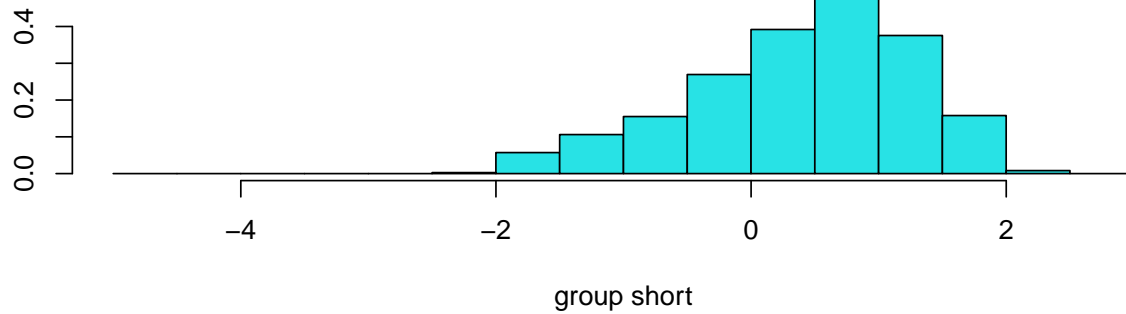
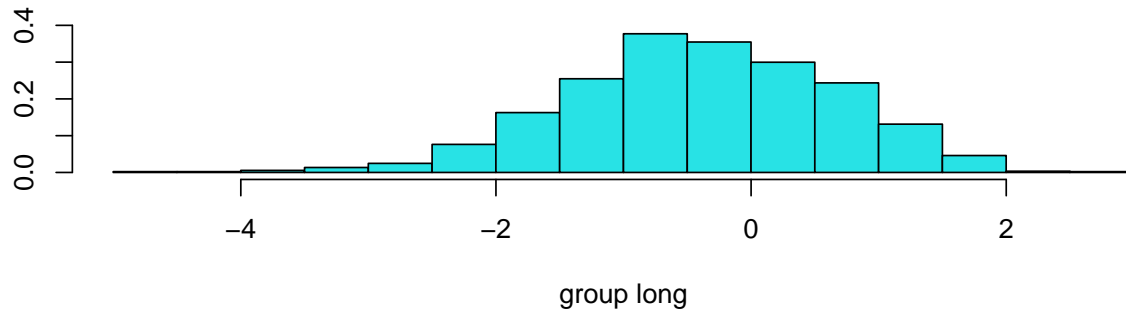
```
## [1] "The train error for MARS is 0.265"
```

```
#test error
test.pred.mars = predict(mars.fit , newdata = x_test_class)
test.error.mars = 1-sum(test_class$recovery_time == test.pred.mars)/length(test_class$recovery_time)
sprintf("The test error for MARS is %.3f", test.error.mars)
```

```
## [1] "The test error for MARS is 0.280"
```

## Linear discriminant analysis (LDA)

```
set.seed(8)
# Using MASS
lda.fit <- lda(recovery_time~.,
               data = train_class)
plot(lda.fit)
```



```
lda.fit$scaling
```

```
##          LD1
## age      -0.005230071
## gender1   0.582852115
## race2     0.111644453
## race3    -0.015332625
## race4     0.078103928
## smoking1 -0.395459244
## smoking2 -0.628005874
## height   -0.497778576
## weight    0.538847538
## bmi      -1.636992668
## hypertension1 -0.368808467
## diabetes1 -0.053193708
## SBP       0.010681574
## LDL      -0.003381055
## vaccine1  0.854566384
## severity1 -0.726871127
## studyB    1.410574304
## studyC   -0.044449352
```



```

set.seed(8)
# Using caret
lda_model <- train(x_train_class, #train x
                  y_train_class, #train y
                  method = "lda",
                  metric = "ROC",
                  trControl = ctrl_class)
summary(lda_model)

##           Length Class      Mode
## prior           2    -none-   numeric
## counts           2    -none-   numeric
## means          36    -none-   numeric
## scaling         18    -none-   numeric
## lev             2    -none-   character
## svd              1    -none-   numeric
## N                1    -none-   numeric
## call            3    -none-   call
## xNames          18    -none-   character
## problemType     1    -none-   character
## tuneValue        1  data.frame list
## obsLevels        2    -none-   character
## param            0    -none-   list

#train error
train.pred.lda <- predict(lda_model, newdata = x_train_class)
train.error.lda = 1 - sum(train_class$recovery_time == train.pred.lda)/length(train_data$recovery_time)
sprintf("The train error for LDA is %.3f", train.error.lda)

## [1] "The train error for LDA is 0.277"

#test error
test.pred.lda <- predict(lda_model, newdata = x_test_class)
test.error.lda = 1 - sum(test_class$recovery_time == test.pred.lda)/length(test_class$recovery_time)
sprintf("The test error for LDA is %.3f", test.error.lda)

## [1] "The test error for LDA is 0.287"

```

## Random Forest - Classification

```

cl <- makePSOCKcluster(no_cores)
registerDoParallel(cl)
rf.grid <- expand.grid(mtry = 1:18,
                     splitrule = "gini",
                     min.node.size = seq(from = 1, to = 50, by = 5))

set.seed(8)
rf_class_model <- train(x_train_class,
                      y_train_class,
                      method = "ranger",
                      importance = "permutation",
                      tuneGrid = rf.grid,
                      metric = "ROC",
                      trControl = ctrl_class)

stopCluster(cl)

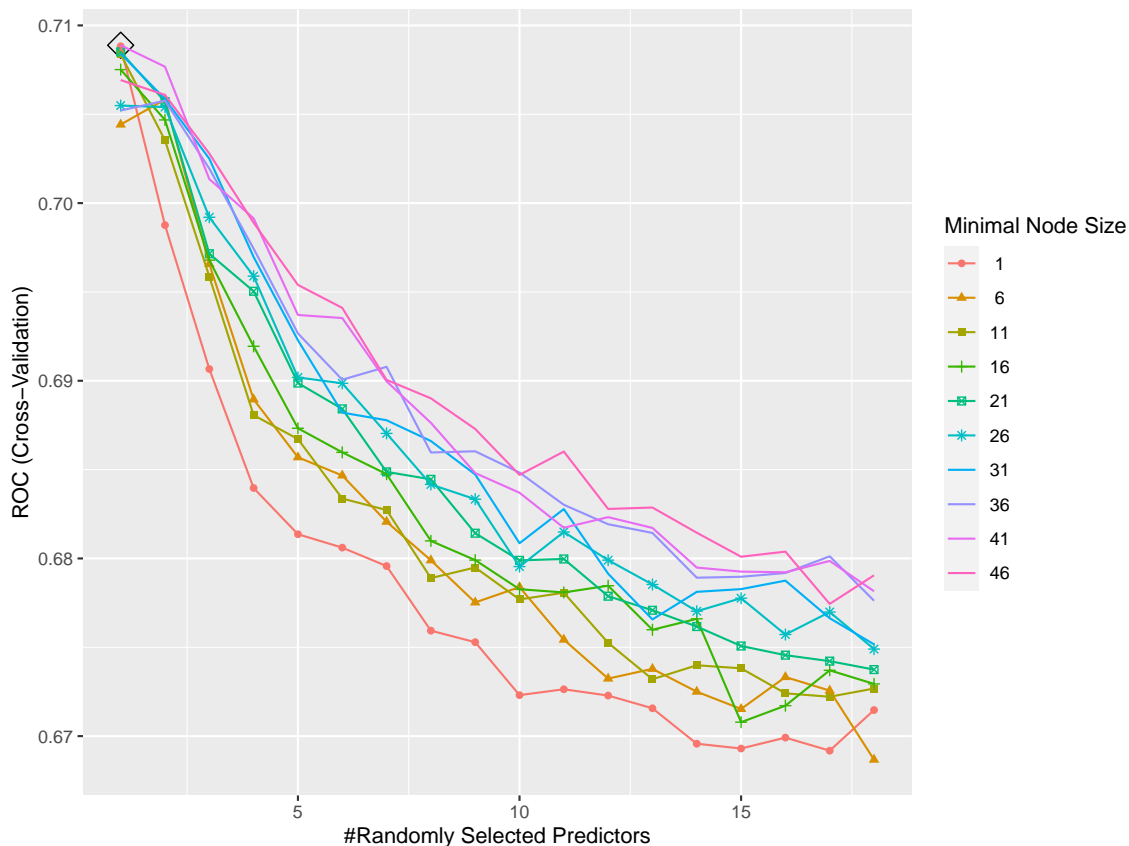
```

```
registerDoSEQ()
```

```
ggplot(rf_class_model, highlight = TRUE)
```

```
## Warning: The shape palette can deal with a maximum of 6 discrete values because
## more than 6 becomes difficult to discriminate; you have 10. Consider
## specifying shapes manually if you must have them.
```

```
## Warning: Removed 72 rows containing missing values (`geom_point()`).
```



```
rf_class_model$bestTune
```

```
## mtry splitrule min.node.size
## 9 1 gini 41
```

```
#train error
```

```
train.pred.rf <- predict(lda_model, newdata = x_train_class)
```

```
train.error.ef = 1 - sum(train_class$recovery_time == train.pred.lda)/length(train_data$recovery_time)
```

```
sprintf("The train error for LDA is %.3f", train.error.lda)
```

```
## [1] "The train error for LDA is 0.277"
```

```
#test error
```

```
test.pred.rf <- predict(rf_class_model, newdata = x_test_class)
```

```
test.error.rf = 1 - sum(test_class$recovery_time == test.pred.rf)/length(test_class$recovery_time)
```

```
sprintf("The test error for Random Forest is %.3f", test.error.rf)
```

```
## [1] "The test error for Random Forest is 0.293"
```

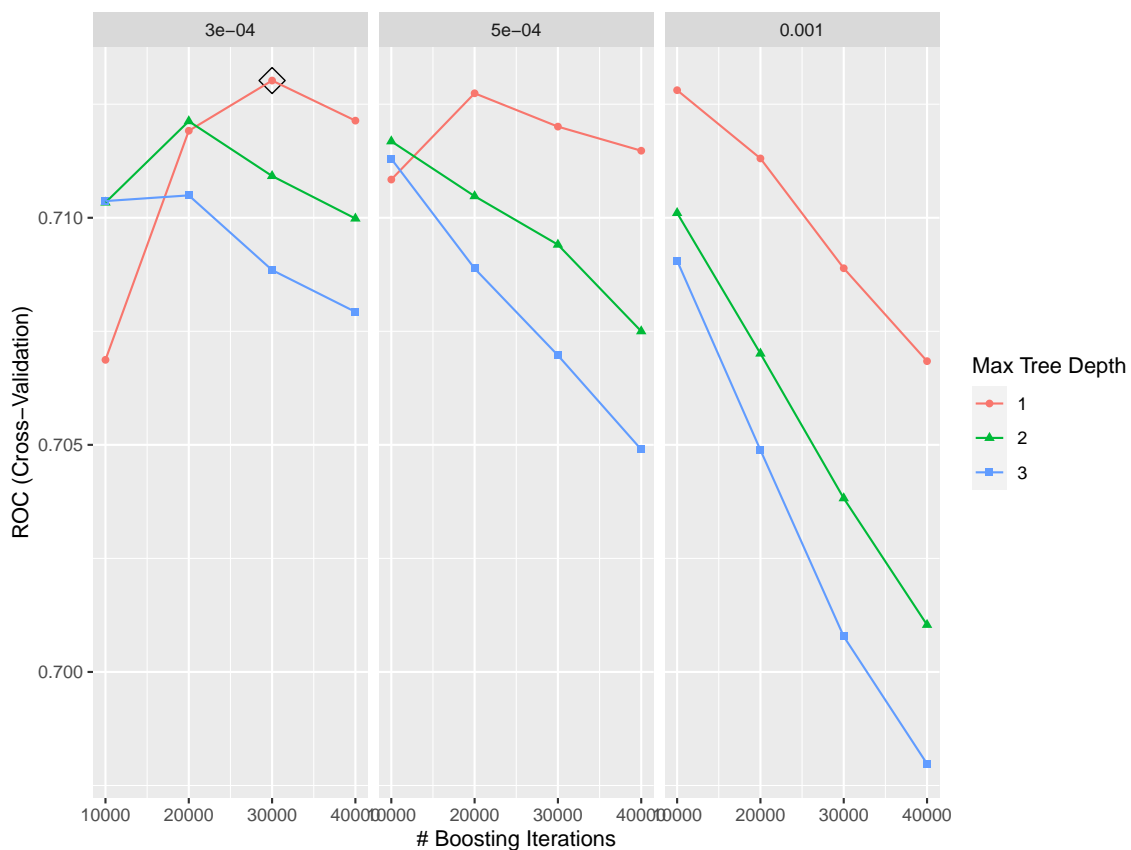
## Boosting

```
gbmA.grid = expand.grid(n.trees = c(10000,20000,30000,40000),
                        interaction.depth = 1:3,
                        shrinkage = c(0.0003,0.0005,0.001),
                        n.minobsinnode = 1)
```

```
cl <- makePSOCKcluster(no_cores)
registerDoParallel(cl)
set.seed(8)
gbmA.fit = train(recovery_time~. ,
                 train_class,
                 tuneGrid = gbmA.grid,
                 trControl = ctrl_class,
                 method = "gbm",
                 distribution = "adaboost",
                 metric = "ROC",
                 verbose = FALSE)
```

```
stopCluster(cl)
registerDoSEQ()
```

```
ggplot(gbmA.fit, highlight = TRUE)
```



```
#train error
```

```
train.pred.boosting = predict(gbmA.fit , newdata = train_class, type = "raw")
train.error.boosting = 1-sum(train_class$recovery_time == train.pred.boosting)/length(train_data$recovery_time)
```

```
# sprintf("The train error for boosting is %.3f", train.pred.boosting)

#test error
test.pred.boosting = predict(gbmA.fit , newdata = test_class, type = "raw")
test.error.boosting = 1-sum(test_class$recovery_time == test.pred.boosting)/length(test_class$recovery_time)
# sprintf("The test error for boosting is %.3f", train.pred.boosting)
```

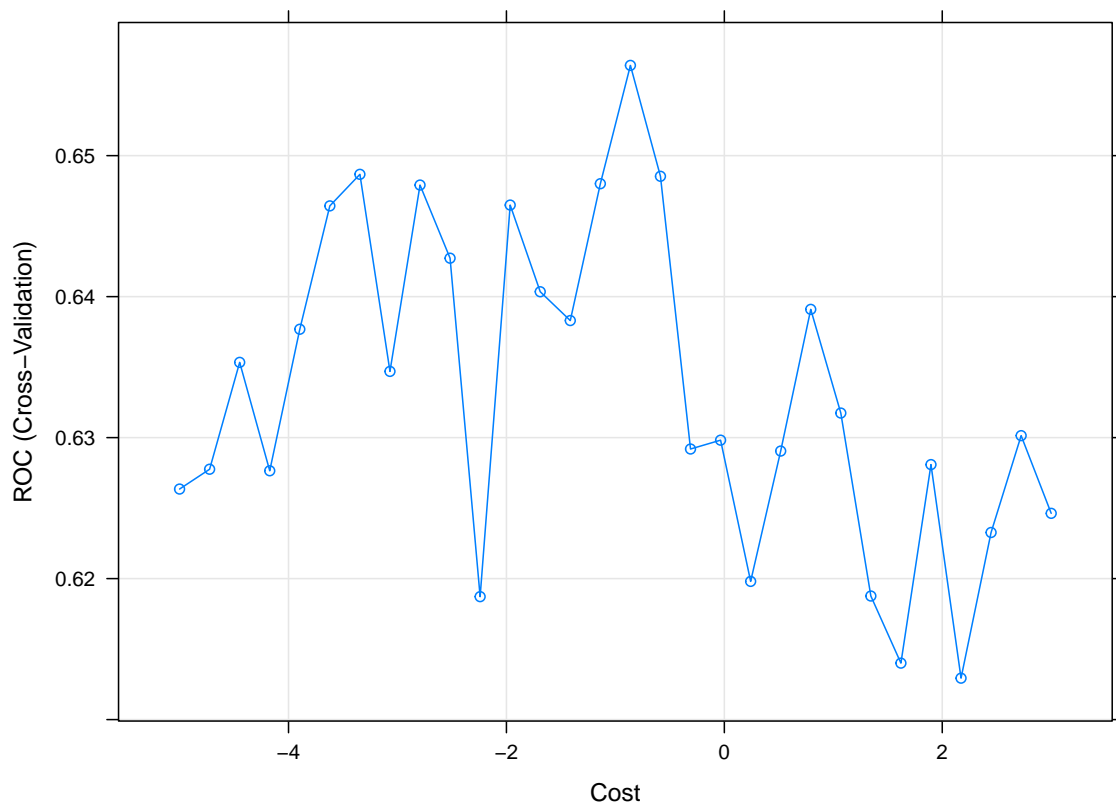
## linear SVM

```
# parallel computing
no_cores <- detectCores() - 1
cl <- makePSOCKcluster(no_cores)
registerDoParallel(cl)

# kernlab
set.seed(8)
svml.fit <- train(recovery_time ~ . ,
                  data = train_class,
                  method = "svmLinear",
                  tuneGrid = data.frame(C = exp(seq(-5,3,len=30))),
                  trControl = ctrl_class)
```

```
## maximum number of iterations reached 0.008313212 0.00758091
```

```
plot(svml.fit, highlight = TRUE, xTrans = log)
```



```
stopCluster(cl)
registerDoSEQ()
```

```
svml.fit$bestTune
```

```
##           C  
## 16 0.4222875
```

```
#train error
```

```
train.pred.svml <- predict(svml.fit, newdata = train_class)  
train.error.svml = 1-sum(train_class$recovery_time == train.pred.svml)/length(train_data$recovery_time)  
sprintf("The train error for Support vector classifier is %.3f", train.error.svml)
```

```
## [1] "The train error for Support vector classifier is 0.292"
```

```
#test error
```

```
test.pred.svml <- predict(svml.fit, newdata = test_class)  
test.error.svml = 1-sum(test_class$recovery_time == test.pred.svml)/length(test_class$recovery_time)  
sprintf("The test error for Support vector classifier is %.3f", test.error.svml)
```

```
## [1] "The test error for Support vector classifier is 0.292"
```

## SVM with radial kernel

```
svmr.grid <- expand.grid(C = exp(seq(-1,5,len=20)),  
                        sigma = exp(seq(-10,-2,len=10)))
```

```
# parallel computing
```

```
no_cores <- detectCores() - 1  
cl <- makePSOCKcluster(no_cores)  
registerDoParallel(cl)
```

```
# tunes over both cost and sigma
```

```
set.seed(8)  
svmr.fit <- train(recovery_time ~ . ,  
                  data = train_class,  
                  method = "svmRadialSigma",  
                  tuneGrid = svmr.grid,  
                  trControl = ctrl_class)
```

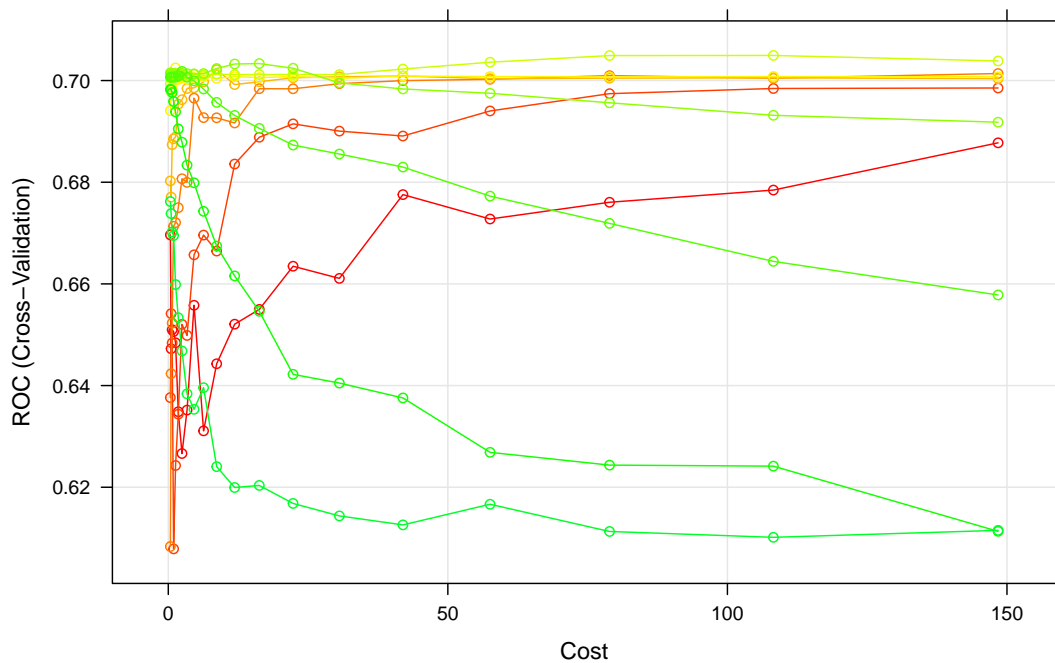
```
## Warning in train.default(x, y, weights = w, ...): The metric "Accuracy" was not  
## in the result set. ROC will be used instead.
```

```
## Warning in nominalTrainWorkflow(x = x, y = y, wts = weights, info = trainInfo,  
## : There were missing values in resampled performance measures.
```

```
stopCluster(cl)  
registerDoSEQ()
```

```
myCol<- rainbow(25)  
myPar <- list(superpose.symbol = list(col = myCol),  
              superpose.line = list(col = myCol))  
  
plot(svmr.fit, highlight = TRUE, par.settings = myPar)
```

			Sigma		
35	○	—	0.00065339197986738	○	—
2	○	—	0.00158932728345653	○	—
35	○	—	0.00386592013947281	○	—
			0.00940356255149521	○	—
			0.0228734649112389	○	—
			0.0556379982778428	○	—



```
svmr.fit$bestTune
```

```
##          sigma          C
## 186 0.00386592 108.2248
```

```
#train error
```

```
train.pred.svmr <- predict(svmr.fit, newdata = train_class)
train.error.svmr = 1-sum(train_class$recovery_time == train.pred.svmr)/length(train_class$recovery_time)
sprintf("The train error for Support vector Machine with radial kernel is %.3f", train.error.svmr)
```

```
## [1] "The train error for Support vector Machine with radial kernel is 0.253"
```

```
#test error
```

```
test.pred.svmr <- predict(svmr.fit, newdata = test_class)
test.error.svmr = 1-sum(test_class$recovery_time == test.pred.svmr)/length(test_class$recovery_time)
sprintf("The test error for Support vector Machine with radial kernel is %.3f", test.error.svmr)
```

```
## [1] "The test error for Support vector Machine with radial kernel is 0.287"
```

After tuning, the cost in this range achieves local maximum accuracy. The value is not on the boundary, hence the local maximum accuracy is achieved. It takes a long time (aprox 30 min) to run so it may be hard to refine into a smaller grid.

## Resampling Results and Model Selection

```
# parallel computing
no_cores <- detectCores() - 1
cl <- makePSOCKcluster(no_cores)
registerDoParallel(cl)
```

```
set.seed(8)
```

```

resamp_class <- resamples(list(enet_logistic = model.glmn,
                             gam = model.gam,
                             mars = mars.fit,
                             lda = lda_model,
                             rf = rf_class_model,
                             boosting = gbmA.fit,
                             linear_svm = svm1.fit,
                             radial_svm = svmr.fit))

stopCluster(cl)
registerDoSEQ()

summary(resamp_class)

```


```

##
## Call:
## summary.resamples(object = resamp_class)
##
## Models: enet_logistic, gam, mars, lda, rf, boosting, linear_svm, radial_svm
## Number of resamples: 10
##
## ROC
##           Min.   1st Qu.   Median     Mean   3rd Qu.     Max. NA's
## enet_logistic 0.6663845 0.6916323 0.7030456 0.7056743 0.7172502 0.7578955  0
## gam           0.6629983 0.7008386 0.7273357 0.7181658 0.7329373 0.7813544  0
## mars          0.6754859 0.6981644 0.7138398 0.7158095 0.7261717 0.7860613  0
## lda           0.6678467 0.6923371 0.7030800 0.7045736 0.7133253 0.7559217  0
## rf            0.6581385 0.6979566 0.7128826 0.7088842 0.7190473 0.7754327  0
## boosting      0.6636805 0.6886796 0.7195946 0.7130213 0.7257773 0.7868205  0
## linear_svm     0.5973526 0.6325420 0.6513609 0.6564019 0.6793198 0.7267970  0
## radial_svm     0.6497614 0.6808119 0.7153345 0.7049381 0.7314020 0.7458827  0
##
## Sens
##           Min.   1st Qu.   Median     Mean   3rd Qu.     Max. NA's
## enet_logistic 0.8882682 0.9016854 0.9299793 0.9236991 0.9410112 0.9606742  0
## gam           0.8651685 0.8867303 0.9157303 0.9130532 0.9353933 0.9606742  0
## mars          0.8820225 0.9007046 0.9073034 0.9124663 0.9186649 0.9606742  0
## lda           0.8932584 0.9073034 0.9271703 0.9270667 0.9438202 0.9662921  0
## rf            1.0000000 1.0000000 1.0000000 1.0000000 1.0000000 1.0000000  0
## boosting      0.9325843 0.9455307 0.9662921 0.9629810 0.9719101 0.9943820  0
## linear_svm     0.9831461 0.9943899 1.0000000 0.9955182 1.0000000 1.0000000  0
## radial_svm     0.8882682 0.9035136 0.9269663 0.9338271 0.9592697 0.9887640  0
##
## Spec
##           Min.   1st Qu.   Median     Mean   3rd Qu.     Max.
## enet_logistic 0.17567568 0.18085894 0.2313032 0.217752684 0.24315068 0.27027027
## gam           0.20270270 0.21959459 0.2448167 0.268196964 0.32427805 0.36986301
## mars          0.17567568 0.21695668 0.2584228 0.270936690 0.33214550 0.36986301
## lda           0.12162162 0.17627730 0.2313032 0.211014439 0.24315068 0.27027027
## rf            0.00000000 0.00000000 0.0000000 0.000000000 0.00000000 0.00000000
## boosting      0.06756757 0.08820807 0.1095890 0.117049241 0.13217327 0.18918919
## linear_svm     0.00000000 0.00000000 0.0000000 0.009570529 0.02392632 0.02739726
## radial_svm     0.10958904 0.13231211 0.1891892 0.179563125 0.21283784 0.26027397
##           NA's
## enet_logistic 0

```

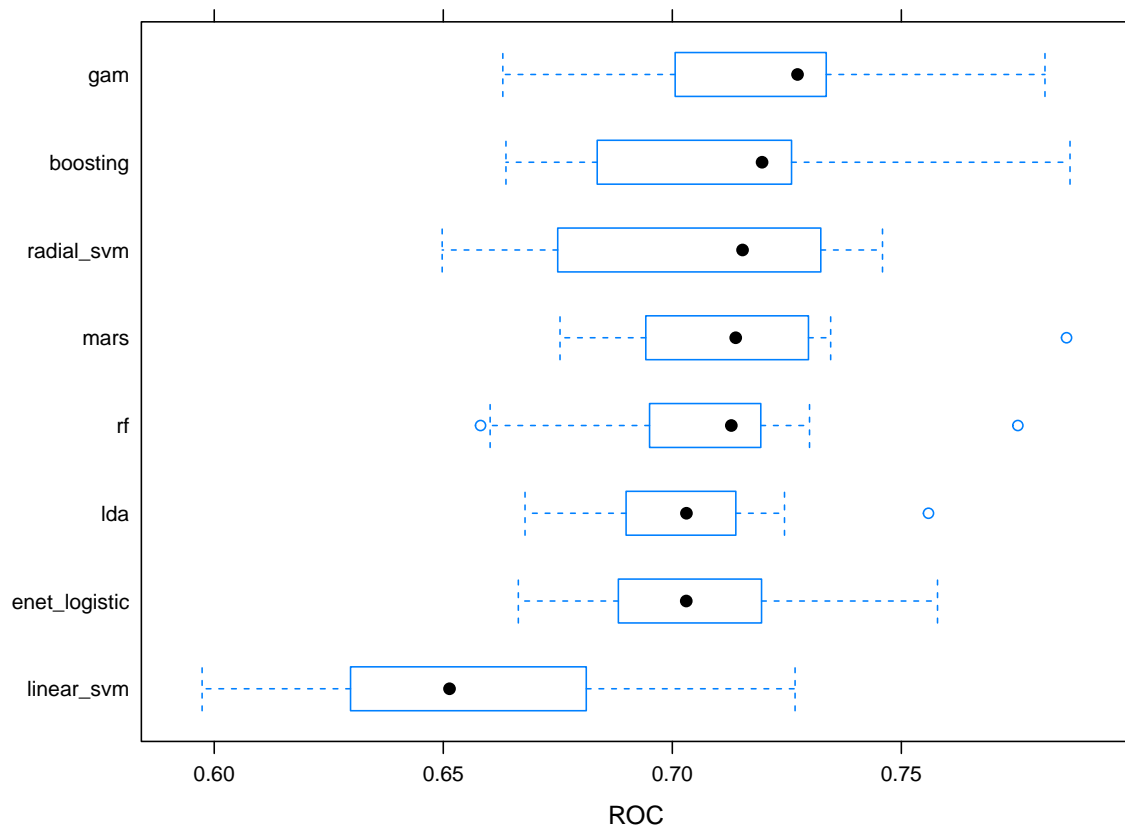
```
## gam          0
## mars         0
## lda          0
## rf           0
## boosting     0
## linear_svm   0
## radial_svm   0
```

```
parallelplot(resamp_class, metric = "ROC")
```



```
bwplot(resamp_class, metric = "ROC")
```






## Final model for classification

```
#feature selection or not
plot(model.gam)
```

GAM



```
#Formula and get the importance of each variable
summary(model.gam)
```

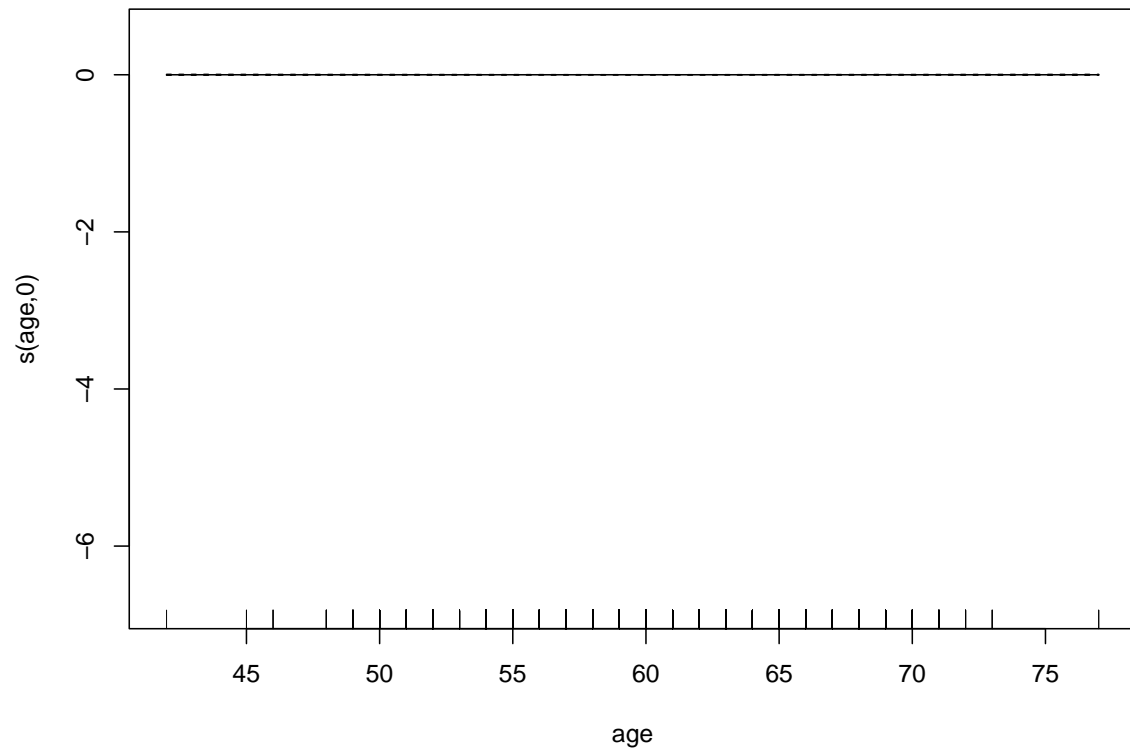
```
##
## Family: binomial
## Link function: logit
##
## Formula:
## .outcome ~ gender1 + race2 + race3 + race4 + smoking1 + smoking2 +
##   hypertension1 + diabetes1 + vaccine1 + severity1 + studyB +
##   studyC + s(age) + s(SBP) + s(LDL) + s(bmi) + s(height) +
##   s(weight)
##
## Parametric coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  -2.07377    0.18672 -11.106  < 2e-16 ***
## gender1       0.47457    0.09509   4.991 6.01e-07 ***
## race2         0.04810    0.21813   0.221 0.825460
## race3        -0.06047    0.12015  -0.503 0.614738
## race4         0.07020    0.16087   0.436 0.662543
## smoking1     -0.33680    0.10828  -3.111 0.001867 **
## smoking2     -0.50352    0.16203  -3.108 0.001886 **
## hypertension1 -0.34676    0.20422  -1.698 0.089510 .
## diabetes1    -0.07480    0.12929  -0.579 0.562881
## vaccine1      0.69695    0.10061   6.927 4.29e-12 ***
## severity1    -0.61790    0.16782  -3.682 0.000232 ***
## studyB        1.18222    0.13684   8.639 < 2e-16 ***
## studyC        0.01115    0.17630   0.063 0.949592
```

```

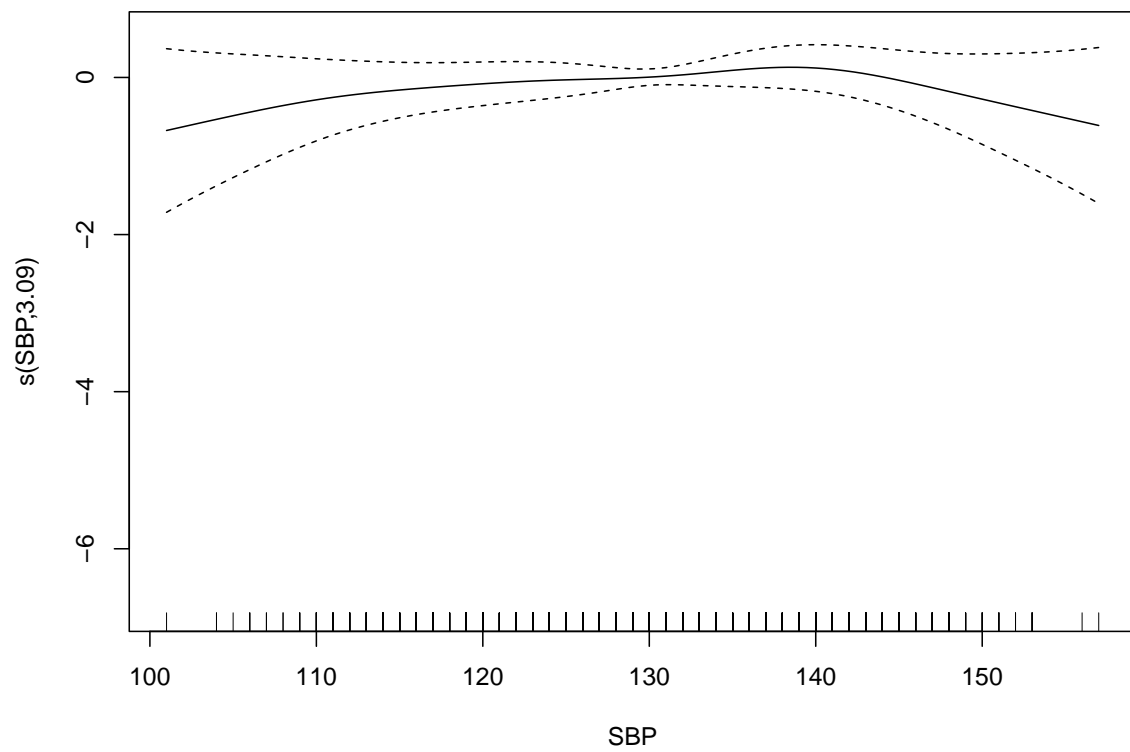
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Approximate significance of smooth terms:
##              edf Ref.df Chi.sq p-value
## s(age)      1.899e-04     9  0.000  0.4443
## s(SBP)      3.090e+00     9  3.726  0.2920
## s(LDL)      1.982e-01     9  0.253  0.2577
## s(bmi)      3.872e+00     9 70.254 <2e-16 ***
## s(height)   1.981e+00     9  5.539  0.0416 *
## s(weight)   9.215e-05     9  0.000  0.9019
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## R-sq.(adj) =  0.131   Deviance explained = 12.4%
## UBRE = 0.076042   Scale est. = 1           n = 2517
#degrees of freedom
model.gam$finalModel

##
## Family: binomial
## Link function: logit
##
## Formula:
## .outcome ~ gender1 + race2 + race3 + race4 + smoking1 + smoking2 +
##      hypertension1 + diabetes1 + vaccine1 + severity1 + studyB +
##      studyC + s(age) + s(SBP) + s(LDL) + s(bmi) + s(height) +
##      s(weight)
##
## Estimated degrees of freedom:
## 0.0002 3.0898 0.1982 3.8717 1.9808 0.0001 total = 22.14
##
## UBRE score: 0.07604226
#discover age and recovery time, df = 0.0002
plot(model.gam$finalModel, select = 1)


```



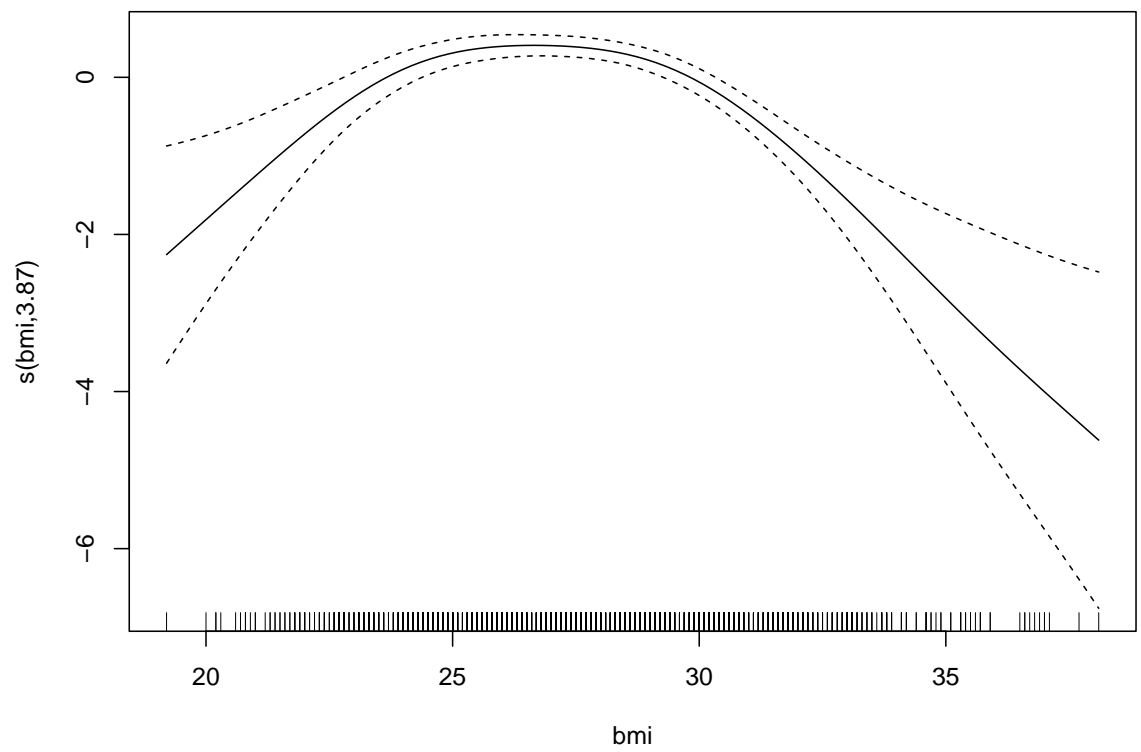
```
#discover sbp and recovery time
plot(model.gam$finalModel, select = 2)
```



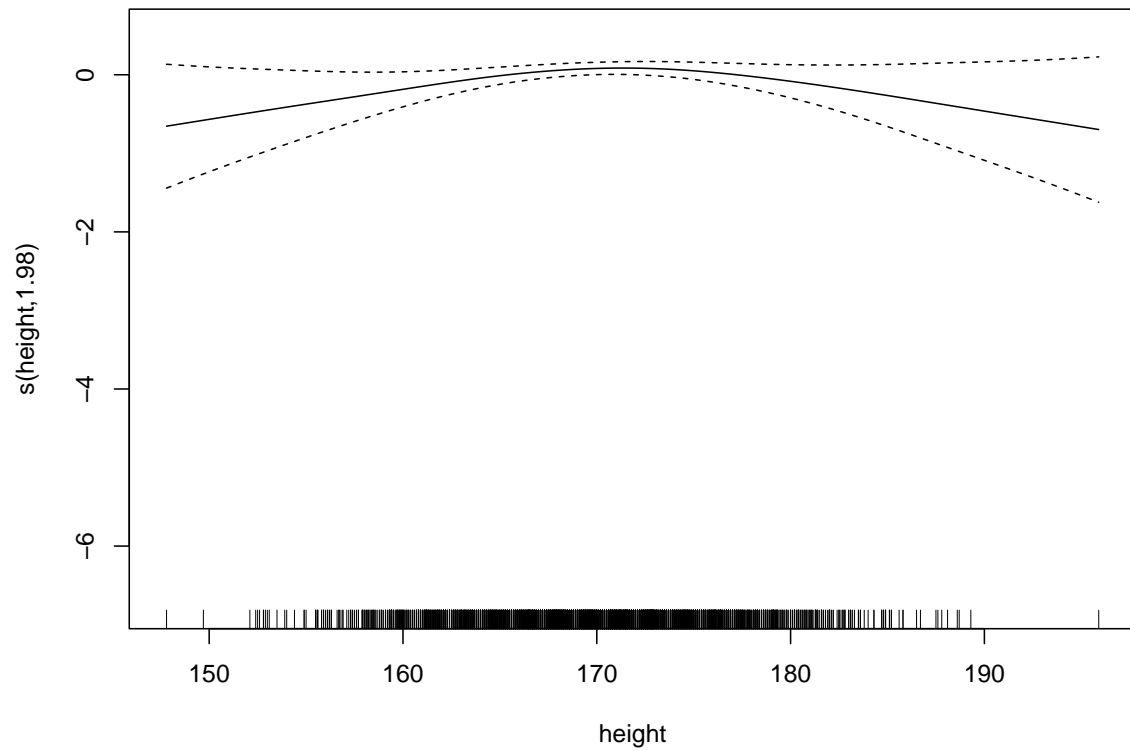
```
#discover ldl and recovery time
plot(model.gam$finalModel, select = 3)
```



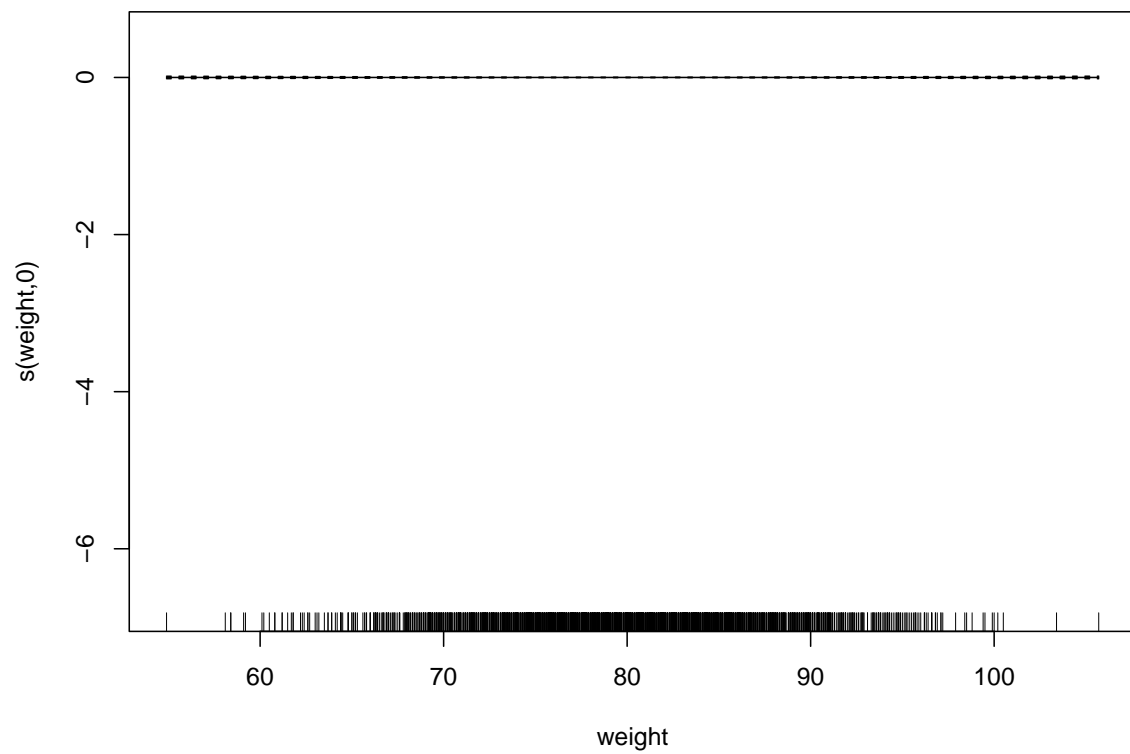
```
#discover bmi and recovery time
plot(model.gam$finalModel, select = 4)
```



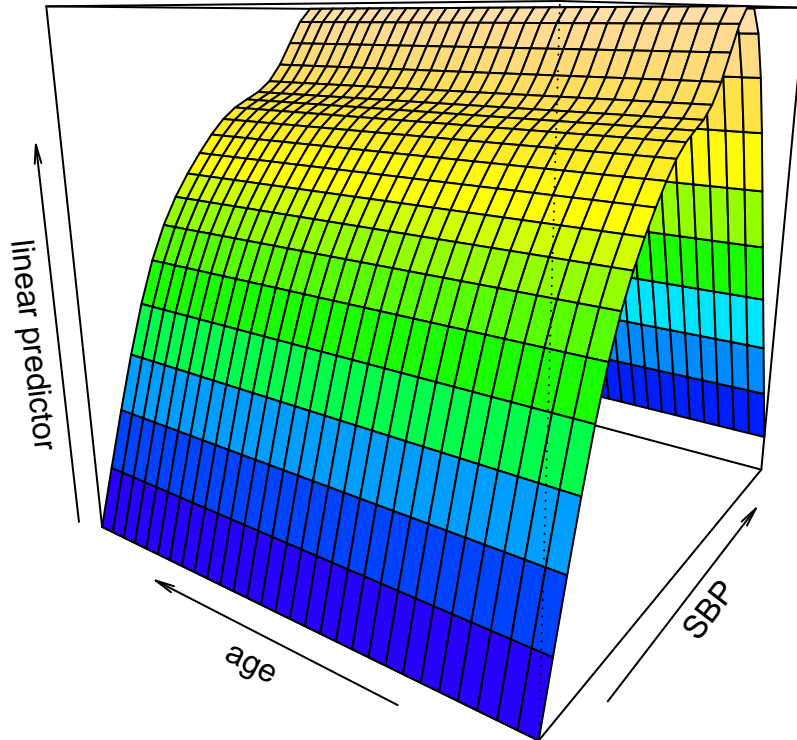
```
#discover height and recovery time
plot(model.gam$finalModel, select = 5)
```



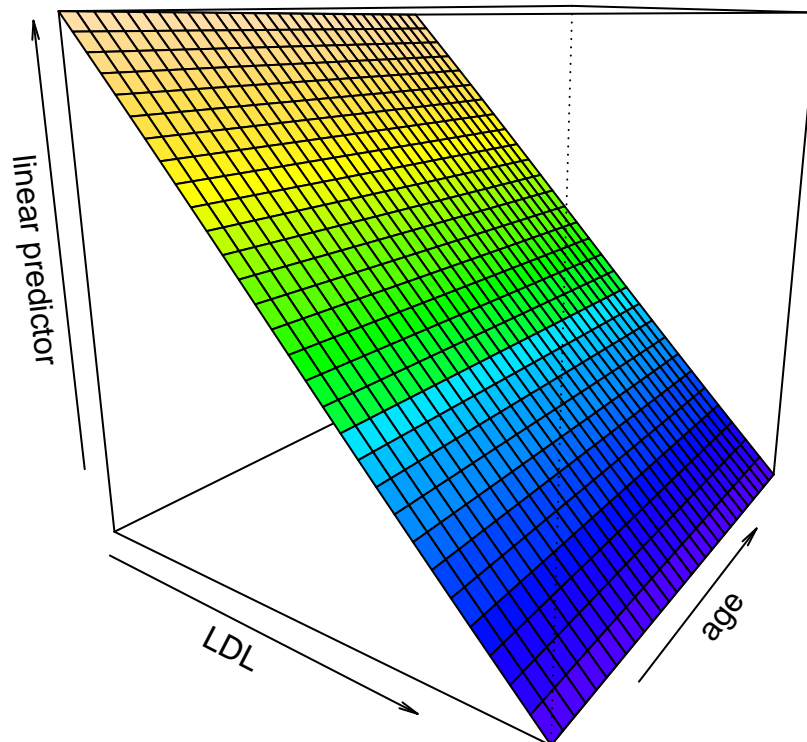
```
#discover weight and recovery time, df = 0.0001
plot(model.gam$finalModel, select = 6)
```



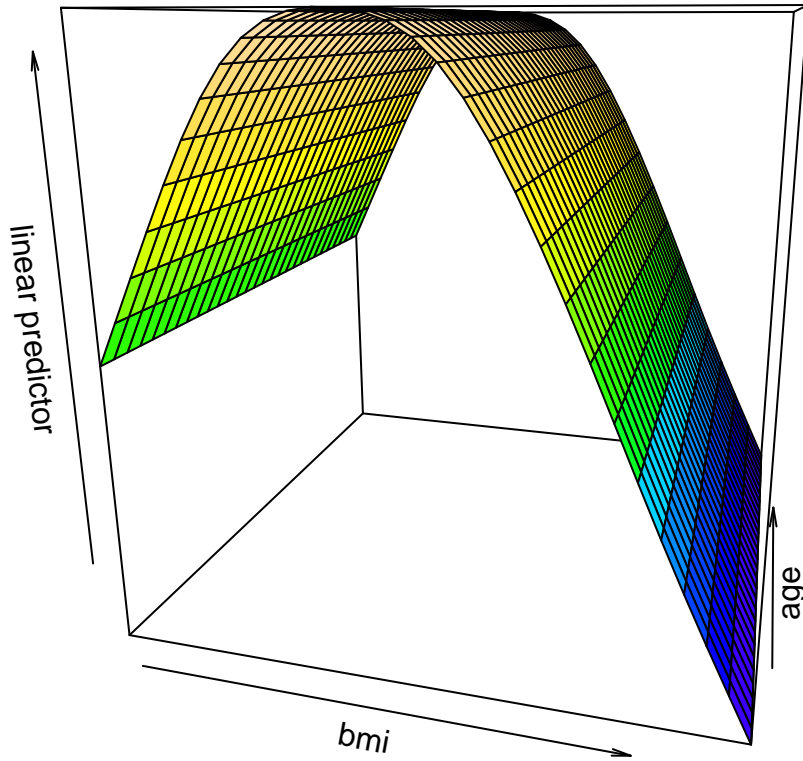
```
#can also visualize the pairwise relation:
vis.gam(model.gam$finalModel, view=c("SBP", "age"), color = "topo", theta = -55)
```



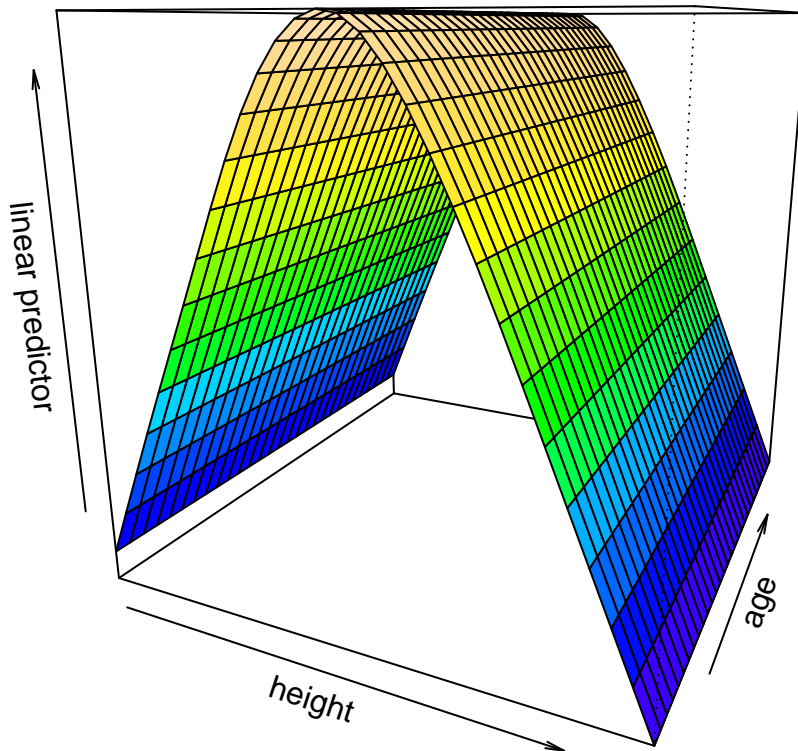
```
vis.gam(model.gam$finalModel, view=c("LDL","age"), color = "topo", theta = 35)
```



```
vis.gam(model.gam$finalModel, view=c("bmi","age"), color = "topo", theta = 15)
```

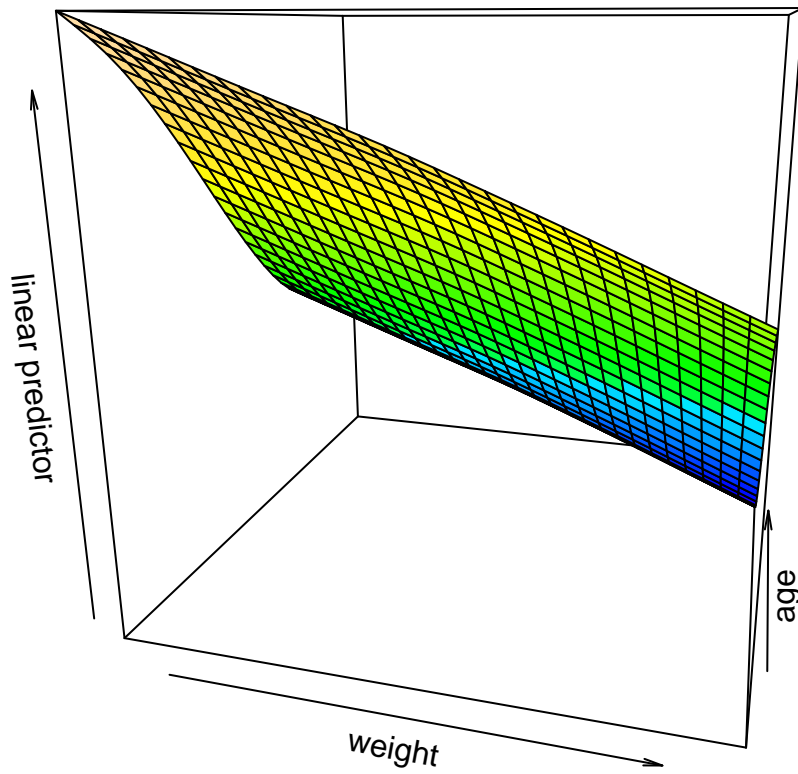


```
vis.gam(model.gam$finalModel, view=c("height","age"), color = "topo", theta = 25)
```

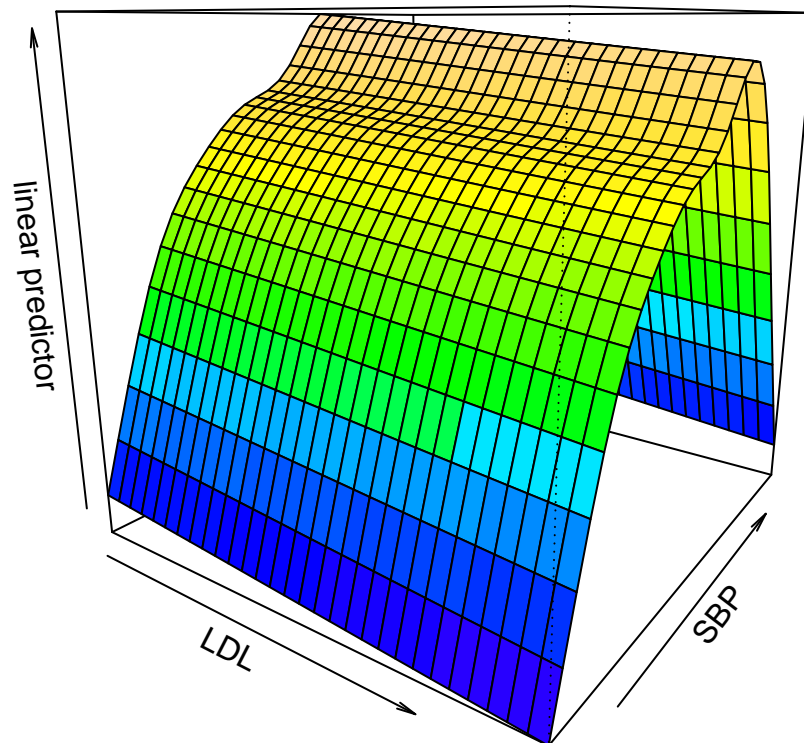


```
vis.gam(model.gam$finalModel, view=c("weight","age"), color = "topo", theta = 15)
```

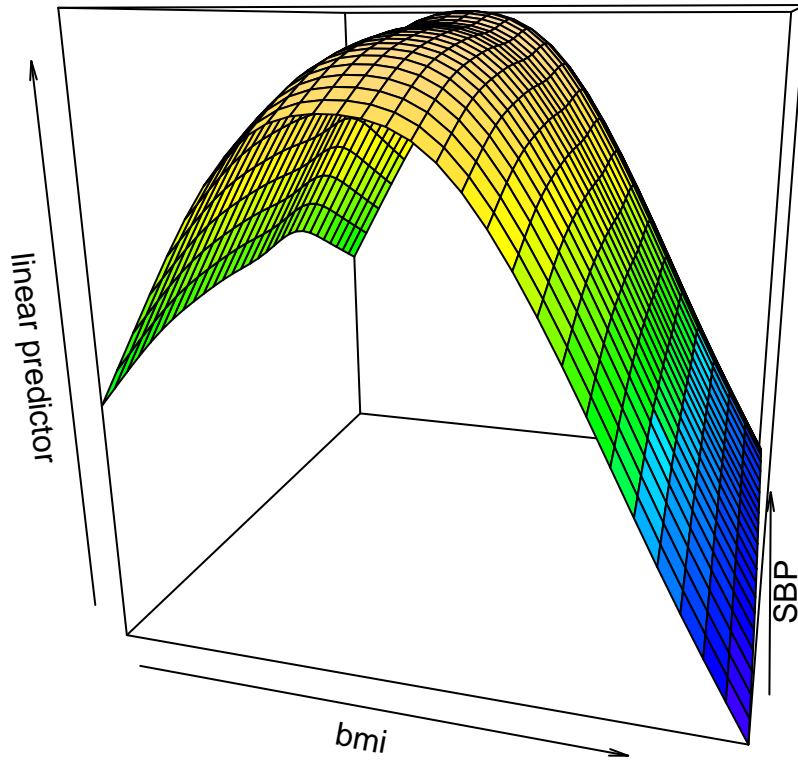




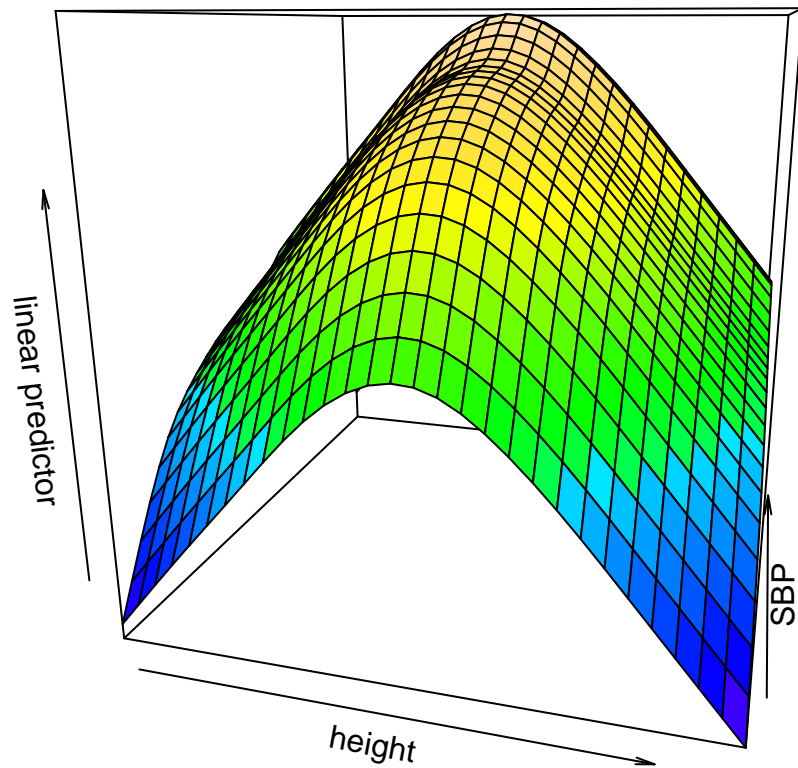
```
vis.gam(model.gam$finalModel, view=c("LDL","SBP"), color = "topo", theta = 35)
```



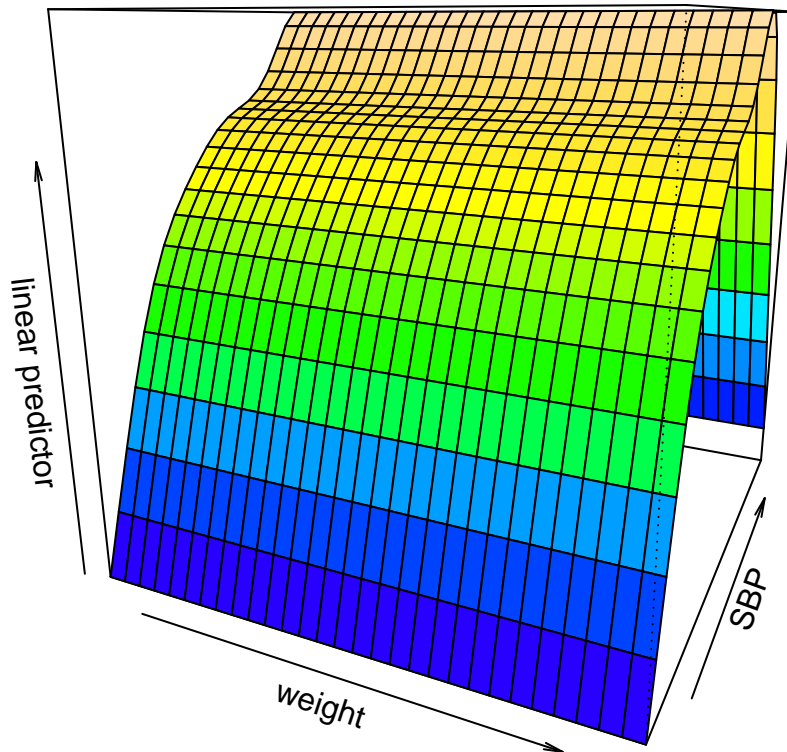
```
vis.gam(model.gam$finalModel, view=c("bmi","SBP"), color = "topo", theta = 15)
```



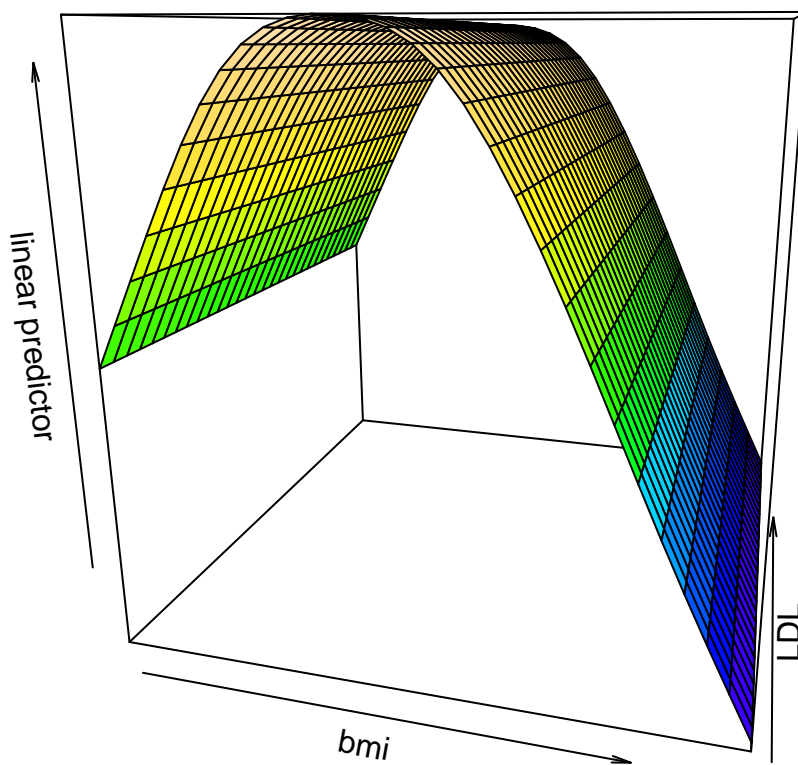
```
vis.gam(model.gam$finalModel, view=c("height","SBP"), color = "topo", theta = 15)
```



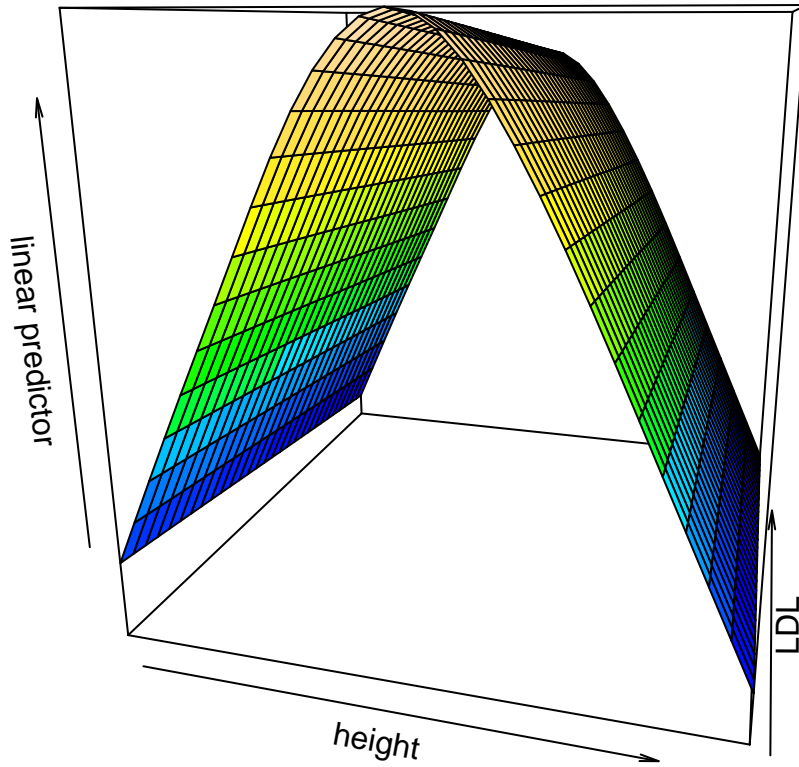
```
vis.gam(model.gam$finalModel, view=c("weight","SBP"), color = "topo", theta = 25)
```



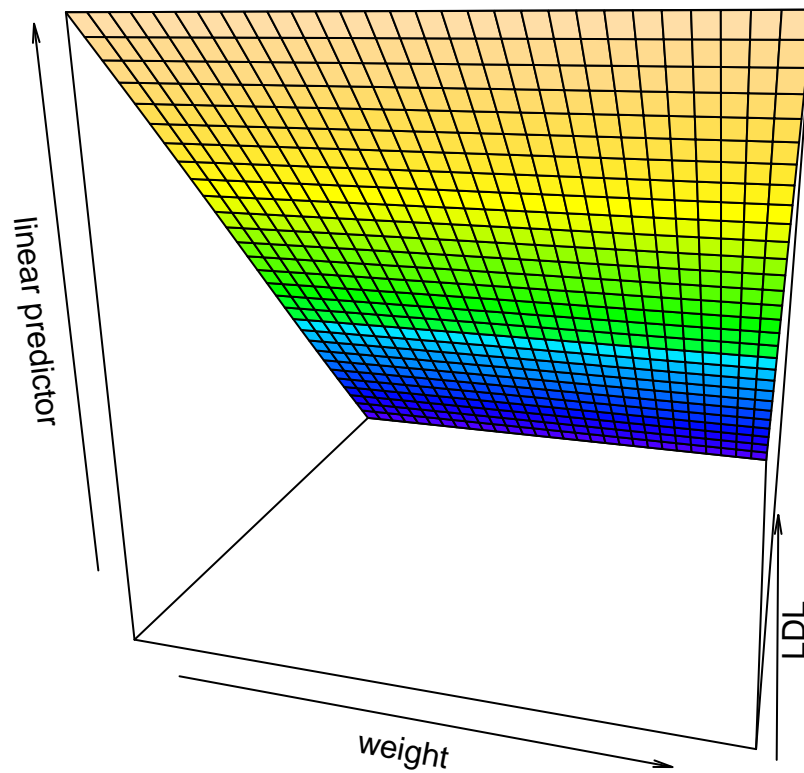
```
vis.gam(model.gam$finalModel, view=c("bmi","LDL"), color = "topo", theta = 15)
```




```
vis.gam(model.gam$finalModel, view=c("height","LDL"), color = "topo", theta = 15)
```




```
vis.gam(model.gam$finalModel, view=c("weight","LDL"), color = "topo", theta = 15)
```



```
vis.gam(model.gam$finalModel, view=c("height","bmi"), color = "topo", theta = 35)
```



```
vis.gam(model.gam$finalModel, view=c("weight","bmi"), color = "topo", theta = 35)
```



```
vis.gam(model.gam$finalModel, view=c("height","weight"), color = "topo", theta = 15)
```

